

[Special Research] Shengshu Vidu Q3 vs ByteDance Seedance 2.0



Questions this study aims to answer:

- Differences between Shengshu Vidu Q3 and ByteDance Seedance 2.0 in terms of video generation effects, speed, etc.
- Technical Dissection Behind the Differences in Generation Effects
- Development Potential and Challenges of Vidu and Seedance from the Evolution Direction of Video Generation Models
- Under the choice of technical roadmap, the differentiation of monetization and future development directions

Brief

Currently, both are in the first echelon, but in the short term, victory or defeat is more determined by data and engineering; starting from the medium term, architecture and training paradigms may push the two towards different endgames: Vidu leans towards physical consistency and exploring the intelligence ceiling of video generation models, while Seedance leans towards controllable productivity and real-time interactive video generation.

- **Capability Comparison**
 - **Strengths of Seedance 2.0: Determinacy and Controllability** — More abundant MultiModal Machine Learning references (including audio), stronger stability in automatic shot planning/narration; post-training (SFT + continuous training + RL multi-dimensional rewards) + inference refiner improve the lower bound of deliverability.
 - **Strengths of Vidu Q3: Physical consistency and cinematic feel** — More focused on the consistency and physical intuition of single-shot continuous motion; TurboDiffusion on the inference side brings speed advantages.
- **Key indicators (current scope, price section to be uniformly verified)**
 - **Generation Duration & Maximum Resolution:** Vidu 1–16s, 1080p; Seedance 4–15s, 1080p/up to 2k.
 - **Input References:** Vidu supports video + image references; Seedance supports combined references of images, text, videos, and audio (up to 12 creatives).

- **Performance and Price:** Vidu has certain advantages in innovative research on inference optimization (TurboDiffusion). This inference efficiency means that when deployed in the cloud, Vidu may have a higher gross profit margin than ByteDance. However, ByteDance also has the ability to optimize across the full stack, from underlying chips to models, and at the same time has multiple application entry points, which means it may be more competitive in terms of price.
- **Monetization Focus:** Currently, the markets and users covered by both parties are similar and converging, but for now, ByteDance's Seedance is focusing its efforts on the C and Pro-C segments, while Vidu may gradually build an advantage in the B segment.
- **Potential points of divergence in the next 12–24 months:** Seedance will assume the role of replacing the existing short video generation workflow (productivity tool --> the foundation of an interactive content platform); Vidu will be more about creative emergence, serving as the "camera" for generating native film and television content & providing realvisual dynamics support
 - **Vidu's Challenge:** Unified Sequence + Long Skip Connection, the improvement of model capabilities relies more on subsequent scaling, and the feasibility of scaling needs to be verified; currently relying on long skip connections to maintain consistency, in the pursuit of fully dynamic spatio-temporal consistency modeling, future efforts will need to "remove the crutch" through large-scale data training.
 - **Seedance's Challenges:** The "architecture tax" brought about by the dual-branch DiT/temporal and spatial decoupling, and whether the upper limit of physical emergence will be widened by the unified sequence approach, i.e., it may be difficult to generate the kind of true cross-modal physical law emergence similar to Sora (lacking a unified world representation space to accommodate these physical general knowledge).

Technical Route Comparison

Conclusion and Discussion

- **At this stage, the benefits brought by data and engineering outweigh the differences in architecture.**
 - The choice of architecture, on the one hand, is more closely aligned with the goals of downstream applications and the characteristics of available data; on the other hand, it is also a compromise on computing power and data efficiency (compared to Sora & Veo).

- Regarding Data drive capability improvement: It is not just a matter of data volume, but more likely high-quality data, as well as the construction of data pipelines (cleaning, annotation quality), and enhancing data diversity through synthetic data (especially how to feed physical laws, such as fluid dynamics and gravity, into models via synthetic data)
 - However, the model architecture determines the utilization rate of data to a certain extent. For example, ByteDance's decoupled DiT, although it may be less physically realistic than Sora, is more efficient in absorbing data on human director's intentions (such as camera movement and transition logic).
 - Meanwhile, the choice of architecture still determines the ceiling of the model, and the subtle differences in architecture remain the key to whether the model can possess causal reasoning capabilities. ---> Possibilities for differentiated development in the next stage
- **In the selection of a series of technical routes such as architecture design and training strategies, both Vidu and Seedance aim to address the same set of issues, with only differences in their approaches (excluding improvements brought about by training data and model scale themselves):**
 - **How to Extend Video Generation Duration?**——Transformer Replaces Traditional U-Net Backbone Network
 - **How to achieve better MultiModal Machine Learning alignment?**——Unified Modeling for MultiModal Machine Learning, Unified Token (Vidu) vs MMDiT + Refiner Model (Seedance) in the Inference Phase
 - **How to improve the consistency of video generation?**——Global Skip Connection (Vidu demonstrates certain architectural advantages) vs. Spatio-temporal decoupling architecture + refiner model (Seedance)
 - **How to balance generation quality with the efficiency of inference and training?** - Systematic optimization (Vidu's current differential advantage lies in Turbodiffusion; Seedance has also carried out a series of optimizations)
 - **How to improve the controllability of video generation?**——MultiModal Machine Learning (all are supported, but Seedance demonstrates certain architectural advantages, and Vidu currently only lacks the audio reference modality)
 - **How to improve the fineness and realism of the image (including understanding of physical common sense)?** —— Vidu mainly relies on architectural advantages, while Seedance mainly relies on refined post-training (SFT + continued training + RL multi-dimensional reward model, refiner model in the inference stage).

- In the next stage, model architectures may lead to a divergence in development directions: Vidu may focus on collaboration with world models and 3D film-grade video generation (amplifying the advantage of a high ceiling for emerging upscaling capabilities); Seedance may focus on real-time interactive 3D video generation (amplifying the advantage of precision and controllability).
 - Breakthrough directions for video generation models in the next stage:
 - Improvements in Generation Duration and Quality: Consistency and Physical Understanding of Long Videos
 - Real-time interactive and precise and controllable generation
 - + 3D Mode
- Bonus Question:
 - In addressing the issues of clipping and the consistency of underlying physical logic, Seedance can currently mitigate this problem (with limitations) through the Refiner model. Vidu will subsequently need to rely on large amounts of high-quality data for training.
 - From the perspective of the next stage of evolution of video generation models, both sides need to **address the engineering balance point between Scaling Law and inference efficiency**. Specifically:
 - **Seedance Needs to Reduce the Burden:** As the current architecture is overly reliant on complex manually designed modules (multi-branch, decoupling layer, post-processing model), ByteDance needs to reduce the communication overhead generated by modal interaction through architecture flattening and lightweighting (such as introducing a more efficient MoE expert routing - already done). Only by shedding the heavy architecture tax can we achieve a generation speed sufficient to support real-time interaction while maintaining precision and controllability, and avoid the system from falling into performance bottlenecks due to over-governance.
 - **Shengshu (Vidu) needs to break the chain** : Shengshu currently relies heavily on the "structural crutch" of long skip connections (Skip Connection) to maintain the consistency of the image. To reach the truly real-time interactive world, Shengshu may need to **remove these hard connections during ultra-large-scale training and shift towards native, dynamic spatio-temporal consistency modeling**. At the same time, it needs to break through the physical limitations of video memory bandwidth and, through a more pure unified sequence architecture, enable the model to spontaneously emerge more stable physical logic without "hard constraints".

Detailed Comparison of Technical Routes

Comparison Dimension	ByteDance Vidu Q3	ByteDance Seedance 2.0
Model Architecture & Key Innovations	U-ViT (U-shaped Vision Transformer) architecture, MultiModal Machine Learning joint generation (refer to the modeling concept of the team UniDiffuser) • All as words, Native MultiModal Machine Learning Fusion Approach: Divide the compressed video into 3D Patches, and take all inputs, including time steps, text conditions, and noisy 3D patches, as tokens for joint modeling. • Long Skip Connections: Similar to U-Net, long skip connections are added between the shallow and deep blocks of the Transformer, directly passing the feature maps from the shallow layers to the deep layers.	It is speculated that the architecture continues the two-stream DiT architecture based on MMDiT + time and space decoupling, and this update also performs sparsification on the model. • Unified Audio-Video Joint Generation Framework Based on MMDiT Architecture (Derived from StableDiffusion3): Treats video streams and audio streams as two independent latent variable branches for joint training, but achieves cross-alignment in the Attention (attention mechanism) layer of each layer. • Decoupling of Time and Space: The spatial layer performs attention aggregation within each frame, while the temporal layer focuses on cross-frame attention computation. Text tokens only participate in cross-modal interaction within the spatial layer. The temporal layer divides each frame into windows, thereby achieving a global receptive field in the temporal dimension. • MoE Architecture (speculated): Vertical division of labor within and between modalities, activating different modality experts based on user prompts.
Logical Deduction of Architecture Selection	U-ViT paper is the first work to prove that Transformer can completely replace the traditional CNN (U-Net) architecture in diffusion models.	MMDiT

- can not only leverage the scalability and long-sequence modeling capabilities of Transformer to break the limitation of video generation duration, but also utilize the low-level detail restoration capabilities similar to those of U-Net.
- **All as words**
Philosophy/Everything can be Token
 - All inputs of the diffusion model (noisy image patches, time step t , conditional information c) are regarded as Tokens,**laying a theoretical foundation for unified modeling of MultiModal Machine Learning.**
 - MultiModal Machine Learning deep alignment breaks the master-slave relationship between different modalities, such as text and images being equally important, enhances the model's ability to understand complex instructions, and alleviates the problem of video logical breakdown.
 - The simple token concatenation method also makes the model more flexible when handling the fusion of text, images, and even **3D data**.
- **Long Skip Connection**{addresses information loss and gradient dispersion in deep networks, which is crucial when processing long video sequences (high resolution, high frame rate)}
- **Time synchronization and semantic consistency across different modalities:**Audio and video evolve in their respective feature spaces and are only aligned at critical attention layers,**allowing different modalities to better preserve their own characteristics.**
- **High controllability and strong decoupling.** Facilitates post hoc fine-tuning (RLHF) of a single modality (e.g., only targeting audio).
 - Multiple MultiModal Machine Learning modes can be added
 - Enhancement of the capabilities of MultiModal Machine Learning combination reference and multi-shot expression
 - **Multi-Lens MM-RoPE:** The architecture allows the model to process image features in the spatial layer and sequence logic in the temporal layer through interleaved MultiModal Machine Learning Position Encoding (MM-RoPE). **The temporal layer is more conducive to learning the ability of shot cutting.**
 - Supports interleaved arrangement of visual tokens and text tokens, and can be extended to the training of multi-shot videos - shots are arranged in the chronological order of actions, and each shot has its own detailed descriptive text.
- **Decoupling time and space to address potential issues arising from the dual-branch architecture**
 - Alleviate computational and memory pressure, and improve the efficiency of training and inference
 - Understanding Time Series

- **Strong spatio-temporal consistency:** Reduces issues such as background flickering or object deformation caused by excessive abstraction of deep features.
- **Detail Preservation & Texture Restoration:** The shallow original features contain high-resolution spatial details, which assist the deep layers through long skip connections for accurate restoration, ensuring that the object contours and background details in each frame of the video are not lost during the decoding process.
- **Enhance visual contrast and dynamic details:** Inject recognizable original information through long skip connections, increase the variance of features, and alleviate the over-smoothing problem (feature homogenization) in deep networks.
- **Training Stability & Model Convergence Efficiency:** Improved gradient flow and low-level feature retention, alleviated the optimization difficulties of deep diffusion Transformer in pixel-level prediction tasks. This enhances training stability, and at the same time accelerates model convergence by shortening the effective learning path, providing a structural foundation for constructing deeper video generation architectures.
- **Massive multi-dimensional data governance + refined post-training** (Continued training + SFT + Video RLHF; Multi-dimensional reward model - basic reward, motion reward, aesthetic reward)
 - Optimization of details and stability of generation
- **MoE: Interpretability and Control, Alleviating Computational and Memory Pressure**
- **Inference Phase: Audio-Visual Joint Refiner Model** — A post-processing diffusion enhancement model used for secondary optimization on top of the results generated by the base model, initializing parameters from the pre-trained base model, and performing independent training, SFT, and RLHF.
 - **Enhance Audio-Visual Consistency:** Perform high-frequency residual correction on the generated Coarse Grain audio-visual stream in the latent space to optimize lip-sync, rhythm alignment, and semantic matching.
 - **Enhanced Local Detail Representation:** Refined high-frequency textures in the image (such as skin texture, water transparency) and audio spatial hierarchy (such as ambient sound, reverb)
 - **Mitigate the long narrative drift problem :** Through the sliding window mechanism, long video segments are segmented for optimization and consistency checks to mitigate narrative drift and local distortion problems, but this does not equate to explicit world state modeling.

- **Improved computational efficiency during the inference phase:** By using local residual prediction instead of global redrawing, the audio-visual generation effect is enhanced with low computational increment.

Future Challenges

- Further improvement in capabilities relies more on scale than on engineering means; as scale further increases, the advantages in training stability, convergence speed, and generation quality brought about by its architecture design may be diluted.
 - As the model further scales, for example, to the 10B level, it is necessary to address the memory and computational pressure brought about by the unified sequence architecture and long skip connections
 - Attention should be paid to the impact of human prior represented by long skip connections on the upper limit of Scale (Skip Connection may limit the further abstraction of features in ultra-deep networks, not as good as the more violent aesthetic DiT route represented by Sora)
 - The issue of semantic dilution that may occur as the video generation duration increases needs to be addressed
- Challenges of sparsification: It is difficult to manually intervene in the model to fix a specific defect, and we can only rely on massive amounts of data to allow experts to learn on their own.
- As further scaling occurs, the complexity of the model architecture (and the human priors it represents) may limit the emergence of cross-modal and physical law capabilities, while increasing the burden of communication overhead and the instability of convergence.
 - Artificially designed bipartite architectures may become a source of information loss and are difficult to generate true cross-modal emergence.
 - may find it difficult to truly understand the deep physical unified field, which may cause it to hit a ceiling when simulating extremely complex cross-modal physical laws (such as the deep coupling of fluid mechanics and sound waves).
 - The communication overhead burden increased by the interaction of different modalities at each layer will crowd out computing power, making it less likely to leverage standard Transformer optimizers to squeeze computing power compared to the unified sequence architecture.
 - Attention should be paid to the convergence instability caused by architectural complexity, and the co-evolution of weights for different modalities needs to be maintained

	• Controllability is inferior to Seedance	
Model Potential	Physical Simulation: More inclined towards Sora's approach, with the emergence ability of physical laws potentially higher than that of the dual-branch DiT model ◦ Native consistency: capable of spontaneously emerging cross-modal physical intuition (such as the sense of impact with synchronized sound and image). ◦ Scaling's high ceiling: perfectly adapted to linear attention and standard optimizers. It has certain advantages in computing power utilization and model upper limit. ◦ Cinematic Sense of Visual Effects : ■ By relying on Skip Connection to lock the physical structure, the generated images have strong geometric stability and are closer to real shooting rather than pixel synthesis. ■ When dealing with complex physical phenomena such as fluids (water, fire, smoke) that are difficult to define with rules, the native architecture of Shengshu is usually more natural than ByteDance. • Inference Efficiency : TurboDiffusion Strong Support	Ultimate controllability: Performs stably in fields such as short dramas, marketing, and advertising that require precise manipulation. ◦ When the scale expands, Text Account can carry more complex semantic logic. ◦ The "Decoupled Optimization" (Modality Independence) of modal balance allows developers to set different Scaling steps for different modalities . ■ Predictable audio-visual synchronization performance: Relying on Joint Attention (Joint Attention) , the interaction between modalities is structured. As the scale expands, it is less likely to suddenly generate uncontrollable "sensory chaos" at a certain node like the unified Token architecture, but steadily improves the accuracy of audio-visual synchronization. ■ Stability of commercial delivery: Relying on the cascading system of "Base + Refiner", physical logic errors can be mitigated through later patches. • Inference Efficiency: Leveraging the MoE architecture, long video generation and real-time interaction are achieved at a lower cost. • 3D Modal Fusion: Supported, but more inclined towards real-time interaction and precise object control to connect to AR/VR interfaces or game engines on mobile devices.

	• 3D Modal Fusion: Supported, but more inclined towards multi-view consistency, evolving towards 3D video generation; meanwhile, if aiming to catch up with the Sora route, the fusion of 3D modalities is the key to constructing a physics engine-style world model	
Foundation work (partial)	U-ViT: 《All are Worth Words: A ViT Backbone for Diffusion Models》 https://arxiv.org/abs/2209.12152 UniDiffuser: 《One Transformer Fits All Distributions in Multi-Modal Diffusion at Scale》 https://arxiv.org/abs/2303.06555 Vidu: 《Vidu: a Highly Consistent, Dynamic and Skilled Text-to-Video Generator with Diffusion Models》 https://arxiv.org/abs/2405.04233 Inference Acceleration: TurboDiffusion employs four core technologies such as hybrid attention acceleration, efficient step distillation, and W8A8 linear layer quantization during the inference phase https://arxiv.org/pdf/2512.16093	MMDiT: 《Scaling Rectified Flow Transformers for High-Resolution Image Synthesis》 https://arxiv.org/pdf/2403.03206 Seedance 1.0技术报告: 《Seedance 1.0: Exploring the Boundaries of Video Generation Models》 https://arxiv.org/abs/2506.09113 Seedance 1.5技术报告: 《Seedance 1.5 pro: A Native Audio-Visual Joint Generation Foundation Mode》 https://arxiv.org/pdf/2512.13507 Inference Acceleration: Multi-stage Distillation + System Optimization Achieves 10× Inference Acceleration.

Overview of the Technological Evolution Process

/	Vidu 1.0	Vidu 2.0	Vidu Q1 (25.7.8)	Vidu Q2 (25.9.25)	Vidu Q2pro(26.1.27)	Vidu Q3(26.1.30)
Key breakthrough	Architectural innovation: Proposed the U-ViT architecture, solving the problem of long video	Generation Speed Optimization: Generate a single-segment video with a length of 4 seconds and	Controllable Video Generation --> Launch Reference Video Generation Function (Image,	• Optimized generation details + improved controllability and consistency	MultiModal Machine Learning Reference Capability: Supports	Audio and Video Synchronization

generation(16s)	a resolution of 512P within 10 seconds, with the generation speed increased by three times compared to Version 1.0 Cost-performance optimization: Calculated based on the unit price per second at 720p resolution, the industry average price is 0.6 yuan/s, while Vidu2.0 is only 0.2587 yuan/s, which is 43% of the industry average price.	Subject Reference) • Main problems addressed: ◦ Uncontrollability of text-to-video; Dependence of image-to-video on storyboards ◦ Video generation consistency issues: consistency of single subject and multiple subjects, supporting up to seven subject inputs	◦ "Micro-expression Generation" serves as the core enhancement scenario, achieving breakthrough progress in extreme facial expression changes, push-pull camera movements, generation speed, and semantic understanding • Video Extension Feature: Free users can generate videos up to 30 seconds long, while paid users can extend them up to 5 minutes.	2 videos + 4 images Reference Editing: 6 Major Video Reference Types (Character, Scene, Special Effect, Expression, Action, Texture) Video Refined Editing: Add, Delete, Modify	
-------------------------------	--	---	--	---	--

				Achieved a 3-fold increase in single-task inference speed compared to the Vidu Q1 reference model;		
--	--	--	--	--	--	--

/	Seedance 1.0	Seedance 1.5 pro	Seedance 2.0	Future
Key Breakthrough	Text-to-Video Model: Spatio-Temporal Decoupled MMDiT Architecture; Multi-Source Video Data + Precise Video Tagging - Multi-source data management combined with accurate and meaningful video subtitles enables comprehensive learning across different scenarios; - An efficient architecture design and training paradigm that supports native multi-frame generation and jointly learns text-to-video and image-to-video tasks;	MultiModal Machine Learning Unified Model, supporting joint audio-video generation: Audio-Video Dual-Branch MMDiT MultiModal Machine Learning Joint Architecture: A unified audio-visual joint generation framework of the MMDiT architecture, which achieves precise coordination of visual and auditory streams in terms of temporal synchronization and semantic consistency through a deep cross-modal information interaction mechanism. Multi-stage Data Pipeline: The team designed a multi-stage data pipeline that balances audio-visual consistency, motion expressiveness, and curriculum-based scheduling. A high-quality and diverse data foundation.	Most likely continues the architecture, training strategy + sparsification processing of 1.5pro Compared to Version 1.5 Video Generation Duration Extension (10s to 15s), MultiModal Machine Learning Reference, Multi-Shot Narrative + Self-Storyboarding Capability Enhancement - Generation quality has been significantly improved	The Seedance series has always hoped to achieve precise restoration of the real world through an extreme and unified algorithm framework.

- The carefully optimized post-training method, leveraging fine-grained supervised fine-tuning and video-specific RLHF with a multi-dimensional reward mechanism, achieves comprehensive performance improvement;
- Through a multi-stage distillation strategy and system-level optimization, a 10-fold improvement in inference speed is achieved. Seedance 1.0 only takes 41.4 seconds (NVIDIA-L20).
- During the inference phase, Diffusion Refiner is used to improve the generation effect while ensuring inference efficiency

- **Post-refinement training optimization:** The team used high-quality audio-visual datasets for supervised fine-tuning (SFT) and introduced the RLHF algorithm specifically tailored for audio-visual scenarios. Specifically, the multi-dimensional reward model effectively enhanced the performance of text-to-video (T2V) and image-to-video (I2V) tasks, comprehensively improving motion quality, visual aesthetics, and audio fidelity.
- **Efficient Inference Acceleration:** Further optimized the multi-stage distillation framework, significantly reducing the number of function evaluations (NFE) required for generation. By integrating inference infrastructure optimizations such as quantization and parallelization, more than 10x end-to-end inference acceleration was achieved while maintaining model performance.
- **Used during the inference stage**Audio-visual joint refiner model

- Higher availability under complex interaction and motion scenarios (motion stability and physical restoration ability --> multi-agent interaction and complex motion scenarios)
- Physical accuracy, fidelity, and controllability are significantly enhanced (Instruction following and consistency performance)
- The model supports 15-second high-quality multi-shot audio-video output and has dual-channel audio capabilities

Comparison of Generation Effects

Conclusion and Discussion

From the perspective of video generation effects, both Vidu and Seedance currently belong to the first-tier players in video generation. With synchronized audio and video output, a video generation time of ≥ 15 seconds, richer shot language, more delicate visual presentation, complex motion and special effects expression, stronger instruction following, and the ability to generate videos based on references, on the one hand, it improves the quality of video generation, and on the other hand, it enhances the controllability of video generation, reduces the randomness for users, and increases the effective output ratio in the production pipeline. This enables both parties to reach a level where commercial delivery is possible in scenarios such as short dramas, animated dramas, and short videos.

- Differentiated Advantages of Seedance 2.0: Automatic Storyboarding and Narrative Stability, Enhanced MultiModal Machine Learning Combination Reference Capability, Sound Effect Enhancement (Environmental Sound).
- Differentiated Advantages of Vidu Q3: Physical Consistency and Realism, Creative Narrative, Sensory Synesthesia, and Character Tension.

Specifically, due to differences in aspects such as underlying architecture, data, and training details, the two parties have certain differences in levels such as consistency and controllability of video generation:

- **Controllability and consistency of video generation:** Both have been improved through MultiModal Machine Learning reference capabilities. However, compared to Vidu Q3, Seedance 2.0 supports richer reference modalities (audio) and has stronger modality combinability. Vidu is more focused on **the consistency strength within a single shot/continuous motion** (3D consistency), while Seedance is more focused on **cross-shot narrative consistency + consistency under instruction constraints**.
- **Narrative Ability :** The notable strengths of Seedance 2.0 are its automatic storyboarding ability and director-level narrative ability. However, according to actual test results, while it has strong instruction-following ability, its narrative may feel industrialized and formulaic. In contrast, the generation results of Vidu Q3 are more randomly inspired and visually dynamic, making it more likely to be eye-catching.

- **MultiModal Machine Learning Alignment (Synchronized Audio and Video):** The advantage of Seedance 2.0 lies in **the high synchronization and precise matching of audio and video**. Through an independent audio expert module, it ensures that specific actions (such as opening a door, hitting a ball) and sound effects are perfectly aligned on the timeline. In contrast, Vidu Q3 tends to showcase **the atmosphere and physical synesthesia (sensory crossover) of audio and video**, focusing on the natural integration of ambient sound and visuals.
- **Understanding of Physical Laws:** The upper limit of Vidu's route may be higher than that of Seedance. At this stage, Seedance can mitigate physical reality logic issues such as clipping through the Refiner Model (post-processing model) and expert constraints, and will need to make a leap towards native spatial perception in the future; Vidu may need to enhance physical stability through Scale (such as higher-order data) in the future.
- **Duration of video generation:** Currently, Vidu Q3 supports direct output of 16s videos, while Seedance 2.0 supports 15s. However, in terms of the video extension function, Seedance has stronger controllability. In the future, when the duration of direct output videos increases, Vidu needs to address the video memory pressure of unified sequences, which may easily lead to semantic dilution and narrative logic deviation, while Seedance needs to solve the problems of narrative diversity and communication overhead.
- **Picture Detail: Seedance 2.0 focuses on high-sharpness textures.** With native 2K rendering capabilities, it clearly shows microscopic details such as skin pores and material textures, making it suitable for commercial advertising needs. **Vidu Q3, on the other hand, focuses on cinematic quality.** Its finesse is reflected in the layering of light and shadow, the realism of reflections, and the natural transition of motion blur, resulting in a higher visual depth of the picture.

In terms of generation speed and price:

- **TurboDiffusion gives Vidu an advantage in inference speed**
- Based on the current API and membership subscription prices, Seedance's pricing is more competitive

Detailed Comparison of Generation Effects

Comparison Dimension	ByteDance Vidu Q3 Pro	ByteDance Seedance2.0
Input Format	Q3 Pro	MultiModal Machine Learning reference creatives: text, images,

	- Images, Text 260127: Reference Generation Launched on Q2 Pro 2 videos (up to 1 video of 5-8s or 2 videos within 5s per subject) + 1-4 images MultiModal Machine Learning input - Supports reference to special effects, expressions, textures, actions, characters, and scenes	videos, audio - Up to 9 images, up to 3 videos (total duration less than or equal to 15s), up to 3 audio files (mp3, total duration less than or equal to 15s), and 12 creatives can be combined for input. - Supports natural language descriptions of actions, effects, camera operations, characters, scenes, and sounds																																												
Output Duration	1-16sOptional(Supports expansion)	4-15sOptional(Supports expansion)																																												
resolution	1080P	1080p, up to 2k																																												
Generation Speed	Personal test: 10s, 1080p, about 7 minutes Official: A 4-second video can usually be generated and completed in about 30 seconds. A 16-second video typically takes about 1-3 minutes, with the specific time affected by factors such as the complexity of the prompt, the level of detail required for the visuals, and hardware performance. Vidu's video generation speed is 10 times faster than OpenAI Sora2 and 2 times faster than Google Veo 3 Fast and Grok-imagine-video. On a single RTX 5090 graphics card, the time to generate a 1080p high-definition video has been compressed from the traditional 900 seconds (15 minutes) to 8 seconds.	Community Test: 10s Video, 20 mins (to be verified) Media: The typical generation time for a 5-second, 1080p audio clip is between 90 seconds and 3 minutes																																												
Price	Vidu Q3 pro API is priced at \$9.60 per minute (approximately 69 yuan per minute) <table><tr><th>能力</th><th>模型版本</th><th>时长</th><th>清晰度</th><th>积分消耗</th><th>错峰积分消耗</th><th>价格 (积分单价0.03125)</th></tr><tr><td>- 图生视频 - 文生视频 - 首尾帧</td><td>viduq3-pro</td><td>1-16秒</td><td>1080p</td><td>每秒32积分</td><td>每秒16积分</td><td>1.00</td></tr><tr><td>- 图生视频 - 文生视频 - 首尾帧</td><td>viduq3-turbo</td><td>1-16秒</td><td>1080p</td><td>每秒16积分</td><td>每秒8积分</td><td>0.50</td></tr><tr><td>参考生</td><td>Q2-Pro</td><td>1-16秒</td><td>1080P</td><td colspan="3">第1秒85积分, 后续每秒+10积分</td></tr></table> User subscription payment: Vidu Q3	能力	模型版本	时长	清晰度	积分消耗	错峰积分消耗	价格 (积分单价0.03125)	- 图生视频 - 文生视频 - 首尾帧	viduq3-pro	1-16秒	1080p	每秒32积分	每秒16积分	1.00	- 图生视频 - 文生视频 - 首尾帧	viduq3-turbo	1-16秒	1080p	每秒16积分	每秒8积分	0.50	参考生	Q2-Pro	1-16秒	1080P	第1秒85积分, 后续每秒+10积分			Dreamina AI 3.0 pro 1 yuan/s <table><tr><th>API</th><th>计费方式</th><th>价格</th><th>并发限额</th></tr><tr><td>即梦AI-视频生成3.0 Pro</td><td>按时长计费</td><td>1 元/秒</td><td> - 免费状态: 并发限制为1 - 付费状态: 并发限制为2 </td></tr><tr><td>即梦AI-视频生成3.0 1080p</td><td>按时长计费</td><td>0.63 元/秒</td><td> - 免费状态: 并发限制为1 - 付费状态: 并发限制为2 </td></tr><tr><td>即梦AI-视频生成3.0 720P</td><td>按时长计费</td><td>0.28 元/秒</td><td> - 免费状态: 并发限制为1 - 付费状态: 并发限制为2 </td></tr></table> User subscription payment: seedance2.0	API	计费方式	价格	并发限额	即梦AI-视频生成3.0 Pro	按时长计费	1 元/秒	- 免费状态: 并发限制为1 - 付费状态: 并发限制为2	即梦AI-视频生成3.0 1080p	按时长计费	0.63 元/秒	- 免费状态: 并发限制为1 - 付费状态: 并发限制为2	即梦AI-视频生成3.0 720P	按时长计费	0.28 元/秒	- 免费状态: 并发限制为1 - 付费状态: 并发限制为2
能力	模型版本	时长	清晰度	积分消耗	错峰积分消耗	价格 (积分单价0.03125)																																								
- 图生视频 - 文生视频 - 首尾帧	viduq3-pro	1-16秒	1080p	每秒32积分	每秒16积分	1.00																																								
- 图生视频 - 文生视频 - 首尾帧	viduq3-turbo	1-16秒	1080p	每秒16积分	每秒8积分	0.50																																								
参考生	Q2-Pro	1-16秒	1080P	第1秒85积分, 后续每秒+10积分																																										
API	计费方式	价格	并发限额																																											
即梦AI-视频生成3.0 Pro	按时长计费	1 元/秒	- 免费状态: 并发限制为1 - 付费状态: 并发限制为2																																											
即梦AI-视频生成3.0 1080p	按时长计费	0.63 元/秒	- 免费状态: 并发限制为1 - 付费状态: 并发限制为2																																											
即梦AI-视频生成3.0 720P	按时长计费	0.28 元/秒	- 免费状态: 并发限制为1 - 付费状态: 并发限制为2																																											

	Annual fee-based subscription, Standard Edition 48 yuan/month, 800 points Professional Edition 179r/month, 4000 credits, 25.5% year-on-year discount Flagship Edition 559r/month, 8000 points, 14.2% year-on-year discount 32 points/s 1080p Based on the standard version, the cost of 1080p video per second is 1.92 yuan	Annual fee-based subscription, Basic Edition 27.4 RMB/month, 55 RMB/month starting from the second year, 1080 points Standard Edition: 79 RMB/month, 158 RMB/month starting from the second year, 4000 points, with a year-on-year discount of 23.4% Professional Edition 216.6/month, 432r/month starting from the second year, 15000 credits, with a year-on-year discount of 43.4% Considering the iteration speed, the cost of the following year is negligible 12 points/s 1080p Based on the Basic Edition, the cost of 1080p video per second is 0.3 yuan
Publicity Focus	Born for "Drama": One-click Video Creation, Synchronous Audio and Video Output Supports 16S synchronized audio and video output, 1080P high-definition picture quality, rich shot language, precise shot cutting, multi-language text rendering, and multi-language output	It's not just MultiModal Machine Learning, but a truly controllable way of creation Stable presentation of complex interactions and movements; MultiModal Machine Learning all-round reference; Comprehensive improvement in instruction following and controllability; Synchronous generation of highly realistic immersive sound effects; Wide compatibility with various production scenarios
Edit Control	260127: Reference editing goes live on Q2 Pro, with natural language controlAddition, deletion, modification (addressing the issue of card drawing: intelligent modification, content editing)	Supports role replacement, deletion, and addition for existing videos
Role Consistency	More focused on the strength of consistency within a single shot/continuous motion (3D consistency)	More emphasis on cross-shot narrative consistency + consistency under instruction constraints.
Original Audio		 (Supports audio reference, and new and richer ambient sound effects have been added)

Storyboard & Camera Movement	can complete the switching between multiple shots such as establishing shot, push-in shot, close-up, and reverse shot within a single segment, and users can control the rhythm and camera movement through natural language prompts .	Self-storyboarding and self-camera operation: Without reference creatives, it also supports automatically planning storyboards and camera operations based on prompt descriptions of the plot , and can also achieve effects such as slow motion and camera shake at places where the plot requires. It supports users to generate a complete segment containing multiple camera switches .
Text control of videos	When generating videos, it can accurately render three languages: Chinese, English, and Japanese. From subtitles, road signs to poster copy, all can be directly written into the prompt.	
Other	Actual test experience shows that Wen Sheng's video effects are stronger than those of Keling and ByteDance	Actual experience shows that the controllability and effectiveness of the reference generation feature are better

Monetization Thinking

Conclusion and Discussion

At present, there are no significant differences in business models, as both cover B2B and B2C users and adopt two types of charging models: API and membership subscription. Seedance's advantage lies in the ByteDance ecosystem, with large B2C traffic and abundant entry points. Shengshu's advantage lies in generation efficiency and the improvement of its own product matrix. Currently, Shengshu's B2B progress is good, especially in overseas market expansion.

In the case where the video generation duration has not exceeded the minute level and the controllability of generation has not significantly improved, it still relies on extracting complete segments from a video for processing. Seedance's multi-shot capability may endow the model with more clipable features in video generation. However, Vidu is not inferior in this regard, as all directly generated videos can be clipped into complete and usable segments.

From the perspective of the evolution direction of model capabilities, Shengshu has the potential to catch up in the future with the exploration by overseas leading teams on the path of video generation models towards becoming global models - similar to the route of Sora. However, in terms of monetization at the current stage, Shengshu's differentiation may focus on **creative** video generation, while being friendly to scenarios with higher requirements for visual realism. For video generation that requires higher controllability, Shengshu can achieve this by encapsulating a system of multi-Agent collaboration, similar to its Vidu Agent launched for marketing scenarios.

Seedance's notable advantage lies in its controllability. In choosing its technical path, it has already internalized the capabilities of many video creation workflows into the model. Therefore, Seedance is more user-friendly for creators who engage in re-creation content and have high requirements for controlling the creative process. Meanwhile, as a member of the ByteDance Seed portfolio, whether in terms of technology roadmap selection or monetization, it seems to place greater emphasis on leveraging the advantages of the existing product matrix to activate the creative passion of C-end users (in the world model, other teams are responsible for research). The future development roadmap of Seedance may gradually evolve from replacing existing editing tools to the stable and controllable generation of movie-grade videos. Additionally, whether it is the decoupling of multimodal display or the decoupling of time and space, it may give Seedance certain advantages in exploring real-time interactive video generation in the future.

Other

At present, the two giants are converging in their business models (covering both B and C ends, with a dual-track system of API + subscription), but they differ in their market entry points and future development directions:

- **Seedance:** Its core competitive edge lies in ecosystem penetration. Leveraging the massive traffic pool and entry advantages of Douyin and Jianying, Seedance has achieved extremely high C-end penetration. Its business certainty stems from the ultimate cost reduction of the existing short video/advertising production pipeline.

- **Vidu:** The core barrier lies in vertical explosive power. Leveraging the native generative efficiency and visual tension brought by the U-ViT architecture, Shengshu has demonstrated strong scalability in the B2B segment (especially in overseas markets and high-premium creative scenarios), and is gradually building its own product matrix Competitive Edge.

Product differentiation focus (not mutually exclusive): Substitute in the existing market vs. Definer of the incremental market

- **Seedance: The Super Post-production that Reshapes Workflows.** Seedance's notable advantage lies in explicit controllability. Its technical approach directly internalizes complex video creation logic (such as storyboarding, local control, and narrative stability) into the model through explicit decoupling of spatio-temporal and MultiModal Machine Learning information.
 - **User Portrait:** More friendly to secondary creators, e-commerce marketers, and **B-side/Pro C users with extremely high requirements for certainty**.
 - **Core Value:** It represents a dimensionality reduction strike against existing editing tools and video production workflows. What Seedance pursues is **an extremely high effective output ratio**.
- **Vidu: A creative engine that reaches the upper limit of art.** Vidu's advantage lies in its native sense of physics and narrative tension. It performs better in physical consistency, lighting realism, and physical synesthesia with audio-visual homology.
 - **User Portrait:** Targeting creators who have extremely high requirements for **aesthetic premium and atmosphere** in areas such as film-level storytelling, comic dramas, and large IP derivatives.
 - **Core Value:** It eliminates the bottleneck in the production process. Through native physical emergence, it makes shots that originally required high-cost special effects easily accessible. Even though it may be slightly inferior in controllability, it can still be **supplemented through engineering with multi-agent collaborative systems such as Vidu Agent**.

Potential Differentiation Directions: Real-time Interactive Video Generation vs. World Model

From the origin of the evolution of model capabilities, the long-term valuation logic of the two has already diverged:

- **Possible Endgame of Seedance: A Real-time Interactive "Video Version of Photoshop"**

As the core application of ByteDance's Seed family, Seedance shoulders the mission of activating the C - end creative dividend. Its decoupled architecture reserves structural advantages for real - time interactive video. In the future, it may evolve into a stable and controllable movie - grade production system that supports real - time feedback, completely blurring the boundary between "generation" and "editing".

- **Possible Endgame of Vidu: The "World Model (Simulation 2.0)" Towards General Intelligence** Vidu is more inclined to pursue the first principles of technology, aligning with Sora's exploration path of the world model. Vidu's long-term value lies not only in generating videos for people to watch, but more importantly, through its native understanding of physical laws, it can be deeply optimized for specific physical simulation needs or high-end film and television requirements. Its future business model may be more inclined to technology licensing or providing underlying physical infrastructure for large IPs.

As a core member of the ByteDance Seed family, Seedance exhibits distinct application-driven characteristics in its technical roadmap and commercial implementation. It is more inclined to leverage ByteDance's extensive product matrix (such as CapCut and Douyin) to activate the creative dividends of C-end users, internalizing model capabilities into productivity tools. Additionally, Seedance's decoupling feature is more user-friendly for B-end users with high requirements for local controllability, such as customized advertising and e-commerce marketing.

In fundamental research, ByteDance has achieved specialization through multi-team collaboration (such as parallel research on world models and development of generative models). The evolution logic of Seedance may be a leap from auxiliary editing tools to movie-level stable and controllable generation. Additionally, its explicit decoupling of multi-modal information and modular design in the spatio-temporal dimension not only ensures the lower limit of commercial delivery but also reserves the structural advantages of computational efficiency and local responsiveness for its exploration of real-time interactive video in the future.

Comparison of current monetization paths

	Number of births Vidu	ByteDance Seedance
	MaaS + SaaS + Agent + APP	

Business
Model

Users and businesses cover content creators, content agencies, and content industry enterprises in over 200 countries and regions worldwide.

AI Video "Subject Community": Lower the threshold for professional video creation; make professional capabilities accessible to everyone; support the sharing, trading, and authorized use of subjects; turn creativity into real assets and achieve continuous monetization

AI Marketing: Vidu Agent, supports multiple languages, multiple voices, multiple models, and multiple scenarios, and on this basis, it supports uploading background music BGM, supports deleting and adding voiceovers, and supports Storyboard editing

Vidu Q2 Pro:Comprehensive Upgrade of Reference Editing: Reference generation supports simultaneous creation with 2 videos + 4 images; 6 major video reference types; video editing enables flexible creation, easily achieving addition, deletion, and modification; better adaptation to multiple scenarios, meeting the creation needs of different fields such as short dramas, animated dramas, and film and television; allowing videos not only to be generated by reference but also to be precisely edited

Vidu Q3: Supports 16s synchronized audio and video output, 1080P high-definition video, with accurate lens control, precise shot switching, and richer audio-visual language; precise rendering of on-screen text makes videos more realistic and natural; supports multi-language output in Chinese, English, and Japanese, enabling zero-threshold global creation; makes the creation of animated dramas, short dramas,

API + APP + Chatbot invocation, fully embedded in ByteDance's model platform and product line

Relying on ByteDance's product matrix: Dreamina, Xiaoyunke, Doubao, CapCut, and Volcano Ark

- Reference images can accurately restore the composition of the picture and the details of the characters
- Reference videos support the replication of shot language, complex action rhythms, and creative special effects
- The video supports smooth extension and connection, can generate continuous shots according to user prompts, and not only generates but also "continues shooting".
- Editing capabilities are also enhanced, supporting role replacement, deletion, and addition to existing videos

	and film and television dramas more professional and efficient	
Potential future monetization directions	• High-premium creative content ◦ Professional film and television-grade company ▪ Deeply optimized for specific physical simulation requirements or high-end film and television needs ▪ Derivation of Big IPs and Digital Asset Simulation • 3D Film and Video Generation & World Model (Gamification and Embodied Intelligence Simulator)	• Deterministic production process ◦ Professional Film Workshop & Content Creator ◦ "Cost reduction and efficiency improvement" in areas such as marketing, e-commerce, and short videos. ◦ Batch production of IP content ◦ C - end subscriptions and traffic monetization for profit • Real-time interactive (including 3D), precisely controllable video generation; integration with terminal devices such as smart glasses

Appendix

 即梦 Seedance 2.0 使用手册（全新多模态创作体验）

Product Measured Reference Materials

Seedance 2.0 Dreamina Effect Evaluation

https://www.bilibili.com/video/BV1ibc5zXEnE?spm_id_from=333.788.videopod.episodes&vd_source=4f5e104198188c2dc2ee48c3b8baf47e

- The video generation from the first frame image has been extended from the original maximum of 10 seconds to 15 seconds (selectable from 4 seconds to 15 seconds)
 - Previous issues: When the time is too long, the entire picture is not very stable, especially after a camera cut, the images of objects in the picture will all change to some extent.
- Single-image direct-output cut-shot video
 - Previous issue: Intense motion scenes are prone to glitching
- Shot movement, scene, character, action migration (new)
 - Can replace images and scenes, and can be used to replicate creative advertisements or camera movements and transition effects

- Dynamic comics are created using the existing video dynamic presentation format (new addition)
- Upload storyboard images to directly generate a video (new)
 - Shot Number, Duration, Shot Type, Camera Movement, Frame Content
 - Previously, one had to first use LLM to write a script, then carefully generate storyboard images based on the script content, select the appropriate images, generate the corresponding short video, and finally add voiceovers, background music, and edit it into a finished video.
- Higher controllability: Video extension function (new)
 - Original video + shot images + prompts, extend video duration
- One-shot
 - Previously, the intelligent multi-frame function was used on JiMeng AI for generation
 - The uploaded reference image can not only appear as a frame, but also as an object or person in it
- Emotion Reference Replica (New)
 - Similar to the past action transfer function, but in addition to actions, it can also replicate the microexpressions of characters
- Enhanced Special Effects

Seedance 2.0 Team Effectiveness Assessment

<https://seedanceai.com/zh/blog/seedance-2-0-review-analysis>

Seedance 2.0 is a MultiModal Machine Learning AI video generation model iterated by ByteDance based on 1.0/1.5 Pro, supporting **text-to-video** and **image-to-video** dual modes. Its core advantage lies in addressing the shortcomings of consistency and narrativity in traditional AI videos, capable of outputting professional videos that meet commercial and film-level standards. Relying on ByteDance's video technology accumulation from TikTok, it excels in visual storytelling and motion dynamics.

1. **Native 1080p movie-grade picture quality:** Broadcast-grade output, with clear details, accurate colors, reduced artifacts, and support for professional depth of field, realistic lighting, and camera-level motion blur;
2. **Advanced Audio-Visual Synchronization:** Natively jointly generates audio and video, including contextual sound effects, multi-language lip-syncing (Mandarin / dialects / Cantonese, etc.), and Spatial Audio, without the need for separate post-production adaptation;

3. **Multi-shot Narrative (Core Revolutionary Feature):** Can split the narrative into multiple smooth shots, maintaining **the triple consistency of character, environment, and narrative** to achieve automated professional editing;
4. **Enhanced Motion Synthesis:** Utilizes V2 motion synthesis technology to generate complex actions and professional camera movements (tracking / jib / Hitchcock zoom) that adhere to physical laws, without dynamic motion artifacts;
5. **Excellent Prompt Adherence:** Accurately parse complex instructions (camera angle, lighting, style, rhythm, etc.), reducing trial-and-error prompt engineering;
6. **Ultra-fast Generation:** 5-10 second segments are rendered within 2 minutes, with the Pro version being even faster and supporting rapid creative iteration;
7. **Style Diversity:** Covers a variety of styles including photorealism, 2D/3D animation, anime, watercolor, vintage film, etc., suitable for diverse creative scenarios.

Seedance 2.0 completes core upgrades based on 1.0 (mid-2025) and 1.5 Pro (end-2025), with core changes including:

1. Resolution: Native 1080p+**2K capability**(Pro version);
2. Audio: From None (1.0) → Basic Generation (1.5 Pro) → **Enhanced Spatial Audio + Full Audio-Generated Video** (Supports uploading audio to generate synchronized video);
3. Speed: 30% faster than 1.0, while 1.5 Pro is only 15% faster;
4. Consistency: Character persistence / multi-shot consistency has been upgraded from "Medium / Good" to **Excellent**, with lip-sync approaching motion capture quality;
5. Duration: Standard 5-12 seconds, **Extended mode can reach several minutes**;
6. Camera Control: From Basic → Advanced → **Professional**;
7. Cost: The cost of 1080p video per 5 seconds has been significantly reduced, with more competitive pricing.

Remaining issues: Underwater complex environment (inconsistent movement of secondary elements), audio-visual synchronization (insufficient fine synchronization between musical instrument performance and musical notes)

SeeDance2.0 Geek Movie Points Out Defects

- Understanding of the physical world remains at the level of "pattern matching" rather than "first principles." This causes the model to expose its shortcomings when dealing with complex or uncommon physical interactions
- In long video creation, details tend to drift and content is difficult to be coherent

- When generating purely realistic content based solely on text without reference, Seedance 2.0 sometimes falls short of two competing products, Sora 2 and Veo 3.1, in terms of the authenticity of the generated characters and the finesse of lighting details, and instead places more emphasis on cinematic expression.

Seedance2.0 Identifies Defects

- The model itself has excellent storyboarding and camera movement capabilities, but it mostly adopts a conservative style, and creating outstanding storyboards and camera movements still requires relatively high-quality prompts.
- The level of detail under complex storyboards is still insufficient.
- Microscopic emotional expressions, such as micro facial expressions, are barely satisfactory.
- The usage rate of long videos is not high.
- For shot scheduling with extremely high information density and extremely fine logical granularity, the model exhibits obvious inertia. In the case of MultiModal Machine Learning reference, the consistency is very low.

Vidu Q3 Effectiveness Evaluation

https://www.bilibili.com/video/BV1he65BoEVj/?spm_id_from=333.337.search-card.all.click&vd_source=4f5e104198188c2dc2ee48c3b8baf47e

<https://zhuanlan.zhihu.com/p/2001737940095742364>

- The first 16-second audio-video direct output model
- Multi-shot automatic switching, smoother camera movement, no need for prompts
- The embedding capabilities of the three major languages, Chinese, Japanese, and English, have been enhanced, text rendering is more precise, and the lighting, shadows, and perspective of text are more natural.
- Previous issues: text rendering is blurry, mouth movement is chaotic or static, and does not match the dialogue
- The characters and scenes referenced by the icon remain consistent during the same generation

- In Q3, instead of insisting on a single continuous shot, the video actively establishes spatial relationships, emotional tension, and action rhythm through the switching of multiple shot types such as panoramic, close-up, and medium shots; sound is also incorporated into the editing structure, and the entry and exit of ambient sounds, action sound effects, and background music will all be coordinated with the changes in the picture that occur with the shots.
- In the case of multi-person conversations, Q3 also demonstrates quite mature understanding and generation capabilities. It can not only accurately distinguish the identities of different speakers but also, based on comprehensive judgments of voice, text, and even visual clues, restore each person's dialogue content, tone strength, and emotional trajectory very accurately.

<https://mp.weixin.qq.com/s/2Ku-FsopJN8WJVNmWwgfew>

Vidu Q3

Highlight: Vidu Q3 is positioned as "Born for Drama", with its core breakthrough being the world's first model to support **16-second direct audio-visual output**, achieving movie-level effects with coherent plot, multi-shot scheduling, and synchronized audio and video in a single generation. It has strong director thinking, supports precise camera panning, tilting, zooming, and tracking, multi-language dialogue lip-syncing, and integrated rendering of on-screen text, significantly lowering the threshold for post-production.

Applicable Scenarios: Highly suitable for creating animated dramas, short dramas, and **film and television demos, especially commercial videos that require long takes with a single continuous shot, complex causal narrative (such as object permanence), and extremely high consistency in character facial close-ups and dialogue performances.**

Kling 3.0

Highlights: Keling 3.0 emphasizes "native creation" and "closed-loop workflow". Its core highlight is the **intelligent storyboard** function, which can automatically schedule shot types based on the script; the first-of-its-kind **"image-to-video + subject reference"** can strongly lock in the characteristics of characters, props, and scenes, addressing the consistency pain point of "face-changing" in AI videos; it also supports native-level high-fidelity text rendering and multi-dialect lip-syncing technology.

Applicable Scenarios:

Suitable for long-form narrative creation (such as serialized animations, movies) that requires strict consistency between roles and props, commercial advertisements that need to accurately display text information, and short dramas with customized requirements for specific dialects or multilingual dubbing.

Seedance 2.0

Highlights: Seedance 2.0 excels in **MultiModal Machine Learning Omnipotent Reference** and **Precise Control**. It breaks the single-input limitation, supports mixed input of images, text, audio, and video, and its biggest highlight is the ability to precisely replicate the camera movements, action rhythms, and special effects of the reference video ("hitting the target precisely"). In addition, it has powerful video editing capabilities, supporting seamless extension, character replacement, background modification, and Pairing Videos To Music Beat.

Applicable Scenarios: Highly suitable for video replication with strict requirements for actions and camera movements (such as dance and fight scenes), creative special effects video production (such as transitions and mind-blowing advertisements), as well as scenarios that require partial modification (face swapping, outfit changing) or creative re-creation of existing video creatives.

<https://mp.weixin.qq.com/s/BD7n1JLS1ZwLH-z90ZveZQ>

"The acting of the video characters generated by Vidu Q3 is very realistic, the lines are relatively natural, and its camera work, including camera movement and shot transitions, is very well controlled. Additionally, Vidu Q3 has a high degree of adherence to prompts; it can bring to life whatever prompt you write, and sometimes the results even exceed your expectations."

Overall, Seedance 2.0 places more emphasis on MultiModal Machine Learning reference capabilities, automatic storyboarding, and flexibility in video editing, making it suitable for creators who have high requirements for controlling the creative process and wish to quickly generate highly polished video content; while Vidu Q3 takes simultaneous audio and video output for long videos, shot scheduling capabilities, and stronger physical consistency as its core advantages, making it more suitable for content production scenarios with strong narrative elements and higher requirements for image quality and physical realism.

Fund Intern

The Vidu Q3 Pro's advantages lie in its fast generation speed, flexible editing control, and high efficiency in simultaneous audio and video output, making it suitable for lightweight content creation that requires quick turnaround, such as commercial advertisements, short videos, and other fields that demand high-speed production and eye-catching results.

Seedance 2.0's advantages lie in its rich MultiModal Machine Learning input, intelligent shot composition and camera movement, and strong character consistency, making it suitable for professional-level creation that requires complex storytelling and precise control.

Regarding video duration: The unique 16-second video output limit of Vidu Q3 reaches the average demand for video duration in the short video field, which can basically cover relevant needs. The lower video duration limit of Seedance 2.0 causes Seedance to get stuck at a rather

awkward threshold for many short video requirements, with relatively poor ability to meet the needs of relevant users.

Cost: The relatively high cost of Vidu Q3 brings certain disadvantages, and it does not have an advantage in the production of long-form videos, which are cost-sensitive and have a longer revenue cycle. Seedance 2.0, on the other hand, is characterized by low cost and ease of use, and is more suitable for debugging and producing minute-level creations when time and cost allow. Its disadvantage in output duration will also be offset as the total video duration increases.

Input: Seedance 2.0's more complete MultiModal Machine Learning input capabilities will lead Vidu Q3 in terms of input. In the short video field, the differences in MultiModal Machine Learning input will not have a significant impact, but in the long video field, MultiModal Machine Learning combined input will bring a richer audio-visual experience.

In terms of generation style : Vidu Q3's creative style is more specialized and innovative, while Seedance 2.0 tends to be more conservative, adapting to the needs of the target audiences of both.

Third-Party Evaluation

Artificial Analysis Evaluation: Q3 Pro Text-to-video Analysis Leads (as of January 2026)

Quality Assessment of Video Generation Model Based on Crowdsourced Preferences in Video Arena (User Real Evaluation)

Q3 Pro Text-to-video analysis (no audio) 榜一, Text-to-video analysis (audio) 榜四, 仅次于 Google veo 3.1 preview、fast preview、grok-imagine-video; image-to- video (no audio) 榜五 (Kling 2.5turbo 1080p 榜二) , image-to- video (audio) 榜四 (Kling 2.6 Pro 榜五)

<https://artificialanalysis.ai/video/leaderboard/image-to-video?audio-output=true>

SuperCLUE Evaluation: Seedance 2.0 Leads

Seedance 2.0 > VE3.1 > Vidu Q3 > Seedance 1.5pro > kling 2.6

Progress of biometric monetization

https://mp.weixin.qq.com/s/YF__CZuyK691nOIP4KKITw

Achieve over 10-fold growth in users and revenue by 2025.

In the film and television (animation drama/short drama/film) industry, Vidu covers over 90% of content providers, tool providers, and production agencies across the entire industry, with

partner clients and partners including Sony Pictures, Tencent Animation, Yuewen Group, CCTV Animation, iQIYI, Jiangxi Film Group, Mango TV, Jiangyou Animation, Juri Lu, etc.

In the internet and intelligent hardware industries, Vidu continues to penetrate leading enterprises in various segmented sectors, with clients including ByteDance, Samsung, Wondershare Technology, TAL Education, Alipay, Honor, etc., and is widely used in content production and product interaction innovation.

In the advertising industry, Vidu has become the preferred choice for leading brand owners, service providers, and platform operators, with partners including JD Group, Alibaba 1688, Amazon, Meituan, Focus Media, BlueFocus, L'Oréal, Anta, etc., driving the large-scale application and stable delivery of AI video in commercial advertising.

In the gaming industry, Vidu has served enterprises such as Lilith, 37 Interactive Entertainment Social, Haiyi, and Tanwan Games, helping to upgrade the efficiency of game content production and interactive experience.

In overseas markets, Vidu has been continuously adopted by global creators and enterprise clients, including Pollo AI, PhotoGrid, OpenArt, Hubx, Fal.ai, Eachlabs, FreePik, GensPark, and numerous overseas creators and institutions, with its application scope covering creative tools, commercial content, and platform-level services. Additionally, Vidu is also used by clients in industries such as education, radio and television, and culture and tourism.

<https://mp.weixin.qq.com/s/F6Q2qDKQatxX3jhuKN9YqA>

- **Internet Domain** : 360 Search, Namu Search, Baidu Search, and Baidu Cloud have all been connected to Vidu; meanwhile, Amazon Web Services Marketplace has been deployed to expand into the global market;
- **Hardware Manufacturing**: All product lines of Lenovo AI PCs are connected; in-depth cooperation with device manufacturers such as AR and VR.
- **Film, Television, and Animation**: Officially announced cooperation with 17k Novel Network, securing the adaptation rights to 7 popular web novels with tens of millions of views; officially announced cooperation with Pule Interactive to produce the world's first AIGC promotional video for the Hollywood film Venom: The Last Dance, which is authorized by a well-known IP and generated through whole-link AI.
- **Cultural Media** : Announced that Shaanxi Cultural Investment Group is building an AI digital cultural tourism platform; announced cooperation with leading advertising and marketing company BlueFocus.

- **College Cooperation:** Jointly promote the practice of China's ultra-high-definition image generation with Beijing Film Academy; sign cooperation terms with Beijing Normal University; launch the "College Students' AI Art Season" targeting over 30 well-known domestic universities including Tsinghua University, Peking University, Beijing Film Academy, and Communication University of China, and specially establish the "Shengshu Vidu AI Chinese Style Creation Special Unit".
- **Overseas Expansion:** Collaborated with US Hollywood animation studio Aura Productions to create the industry's first AI science fiction TV series; reached a partnership with Turkey's leading internet company HubX to develop exclusive AI video generation capabilities for its app; and partnered with AI animation industry's TikTok equivalent PocketAnime, jointly launching a global animation competition.