

[Weekly Briefing] 0420-0426

- 💡 What is most worth mentioning this week is the centralized release of models ---> the current evolutionary directions of architecture innovation, system optimization, etc.
- There is a new company NeoCognition (focused on building micro-world models to enable Agent self-learning) that can be continuously monitored, as well as a new interaction method (Flipbook: the model renders each pixel in real-time, bidding farewell to traditional UI)
- Regarding the progress of the AI market, we can refer to Anthropic's report on 80,000 people. From the matrix of **perceived exposure x perceived substitution risk, we may be able to see the implementation progress and future opportunities**
- Openrouter token usage totaled 22T, with a month-on-month increase of 5%, achieving the first month-on-month growth since half a month ago, but still lower than the peak before consecutive negative growth. Kimi 2.6 1.58T, Claude sonnet 4.6 1.36T, DeepSeek V3.2 1.28T, Claude Opus 4.7 1.15T, Gemini 3 Flash Preview 1.04T

Key Trends and Insights

Signal 1: Model Efficiency Revolution

1. Signal Interpretation

Continuing the discussion from the previous weekly report, before the model capabilities have yet to achieve generational leaps, the direction of architectural innovation has begun to partially converge towards improving efficiency. **The MoE architecture and software-hardware co-optimization are becoming increasingly important**, and the significance of DeepSeek V4 lies in successfully running the entire training on purely domestic chips.

- Extension One - Model Architecture Innovation Still Needs to Consider Simplicity :**
However, due to the complexity of its architecture, although it has left profit space for the inference optimization layer infra company, it may need to develop towards a simpler architecture in the future. (Chapter 0426 Star Achievement Potential Space Sharing [国文字记录：大模型推理成本优化讨论 2026年4月26日](#))

- DeepSeek V4 has made many radical improvements to the model architecture in terms of compressing KV Cache to reduce computational costs, and the model scale has also been extrapolated to 1.6T. It is quite remarkable that it can be stably trained (using domestic GPUs).
- Hybrid KV Cache Management: When there is no concise way to solve a series of problems, we can only increase the complexity. This gives the infrastructure team at least a year of leeway to work on it. Originally, everyone was involved. However, the complexity of the model means that the training process is rather opaque, with many chain reactions. Therefore, we still tend to prefer a relatively simple architecture in the future.
- **Extension 2 - In the long run, the inference optimization space left for Infra layer companies: Supernode optimization, cross-cluster optimization.** The work on single chips is mainly left to emerging chip companies to complete. (Chapter 0426 Star Miracle Latent Space Sharing x [【OP】机械系系友论坛：AI基础设施的探讨|观点分享](#))
 - **Why does cross-cluster become possible?**
 - **Compressed KV Cache parameters at the model architecture level** (whether it's Kimi KDA, DeepSeek V4 HCA, or Ant Lynx model, Xiaomi MIMO), ensuring transmission speed, otherwise the bandwidth requirements for transmission would be too high (which also confirms the need for high-speed interconnection - optical communication). On this basis, heterogeneous devices can be further utilized to achieve an improvement in overall cost-effectiveness.
 - Single data center --> Heterogeneous co-location in the same city --> Computing power network: Development also needs to proceed step by step. Currently, the co-location deployment within the same data center, which can be well implemented, and the cooperation between Qijing and Kimi on perfile as service are still in the early stages.
 - **Cross-cluster for Model Training:** Jiliu Technology has verified the cross-95km effectiveness with China Telecom in 2024, and currently Google has successfully completed the practice - successfully training a 120-billion parameter model on a **2-5 Gbps WAN** across four US regions, which is **more than 20 times faster** than traditional synchronization methods. It supports mixed hardware generations (TPU v6e + TPU v5p).



Kimi Linear (KDA - Kimi Delta Attention): A hybrid linear attention architecture launched by Moonshot. Its core lies in a **3:1 hierarchical ratio** : Every 3 layers use the

efficient KDA linear attention, and only 1 layer uses the global MLA (Multi-Head Latent Attention).

- **Effect:** This design directly reduces the KV Cache requirement by **75%**. KDA is responsible for capturing positional information, and MLA is responsible for global modeling. This "combination of light and heavy" allows the cost-effectiveness of long text reasoning to spill over from within the data center to the cross-node level.

DeepSeek-V4 (HCA & CSA): DeepSeek further evolved its compression strategy in V4, proposing **HCA (Heavy Compression Attention)** and **CSA (Compressed Sparse Attention)**.

- **HCA:** Forcefully merge the KV caches of extremely long sequences into extremely small entries, aiming for the highest transmission efficiency.
- **CSA:** Similar to intensive reading, it preserves sparsity on the basis of compression.
- **Effect:** The KV Cache occupancy of the V4-Flash version is only **7%** of that of the previous generation product.

The "Foundation" in the Era of MultiModal Machine Learning and Agents: Xiaomi MiMo and Ant Lynx

- **Xiaomi MiMo (V2-Omni):** A full-modal foundation for the Agent era. Its logic lies in **the deep binding of perception and action**. When Agents frequently call tools, process up to 10 hours of audio, or handle massive amounts of code, traditional full attention mechanisms can cause memory collapse. **MiMo ensures that the overhead of state switching is low enough when Agents execute complex long-path tasks (such as Code Agent) through underlying modal fusion and state compression.**
- **Ant Lynx (Bailing Series):** Ant's path is more inclined towards the "cost-performance balance" of large-scale commercial scenarios. Lynx, when dealing with ultra-long contexts, utilizes a compression technique similar to a "finite-state machine", and through offline preprocessing and online recalculation of key parts, solves the problem of KV Cache waste caused by "prefix mismatch" in RAG scenarios.

- **Extension 3 - Model routing is also important, and there is still room for differentiation in meeting the application requirements of end users, but in the long run, the opportunities probably do not belong to infra companies, but may be at the Agent OS layer or the application companies themselves (for infra layer services, stability remains the top priority, as well as supporting users' diverse routing strategies):**

- Sakana has released the Fugu small model, specifically designed to dynamically coordinate multiple cutting-edge large models (Gemini, GPT, Claude, etc.) to complete complex tasks. There are two variants: fugu-mini (latency-first) and fugu-ultra (performance-first). It is based on the ICLR 2026 papers Trinity and Conductor. Benchmark: fugu-ultra outperforms Gemini 3.1 high, GPT 5.4 high, and Opus 4.6 max in GPQA Diamond 95.1, LiveCodeBench v6 93.2, and SWE-Pro 54.2. It is compatible with the OpenAI format API and can directly replace existing calls. (0427 Frontier Dynamics)
- Discussions on model routing in multi-agent systems emerged relatively early, such as MasRouter: Learning to Route LLMs for Multi-Agent Systems <https://arxiv.org/pdf/2502.11133> However, large-scale applications have not emerged until now. There are some engineering challenges, such as the API being in canary release and each call being different, like Claude often likes to occasionally reduce intelligence to cut costs (0427 Hongyuan sharing)

2. Model Comparison (Launched the Model Frontier Weekly Report Webpage)

- **OpenAI GPT-5.5, officially called "the most intelligent model to date".**

In programming, Terminal-Bench 2.0 reaches 82.7%, SWE-Bench Pro reaches 58.6%, and evaluations by multiple senior engineers suggest that its inference and autonomy significantly surpass GPT-5.4 and Claude Opus 4.7. In knowledge work, GDPval (Agent tasks for 44 occupations) reaches 84.9%, and Tau2-bench Telecom reaches 98.0%. Key Highlights: The model is co-designed with NVIDIA GB200/GB300 NVL72, with inference speed 20% faster than GPT-5.4 (due to Codex itself optimizing the inference stack); API context window is 1M, 400K in Codex; API pricing is \$5 input/\$30 output per 1M tokens. In terms of security, OpenAI has rated the network and biological capabilities of GPT-5.5 as "High" level, deployed more stringent classifiers, and launched the Trusted Access for Cyber project.

- **DeepSeek V4**

Two versions are provided: V4-Pro (1.6T total parameters/49B activated), whose performance is comparable to top-tier closed-source models and is the state-of-the-art in open-source Agentic Coding; V4-Flash (284B total parameters/13B activated), with inference capabilities close to Pro, focuses on high cost-effectiveness. Both models support 1M context by default, adopting Token-wise compression + DSA (DeepSeek Sparse Attention) innovative attention architecture, significantly reducing the computational and memory costs of long contexts. V4 has integrated mainstream Agent tools such as Claude Code and OpenClaw. The API will be available starting today, and the old models (deepseek-chat/reasoner) will be taken offline on July 24. API pricing: V4-Flash \$0.10/M input, \$0.40/M output; V4-Pro \$2.00/M input, \$8.00/M output.

Related Dynamics: DeepSeek has initiated financing, Tencent and Alibaba are in talks with DeepSeek regarding investment, which marks DeepSeek's first acceptance of external financing. The target financing amount exceeds \$300 million, and its valuation has been adjusted upward from the initial approximately \$10 billion to over \$20 billion (some reports indicate \$25.5 billion) due to market enthusiasm.

• **Moonshot Kimi K2.6**

The long-range programming model for continuous Agent work scenarios behaves more like a human engineer. The model has a MoE architecture with 1T total parameters/32B active parameters, includes 384 experts, 8 routing experts, and 1 shared expert, and adopts the MLA attention mechanism.

Associated Breakthrough: Achieved a 54% throughput improvement in 0419 with the Prefill-as-a-Service cross-data center inference architecture of Zhang Mingxing's team. Zhang Mingxing believes that the linear attention architecture is naturally suitable for optimizing such systems.

● 本周重大发布

高 高置信度

GPT-5.5
OpenAI | 4/23 | 正式名称GPT-5.5

不是 GPT-6，命名本身就是信号——性能提升有意义但非范式级。Terminal-Bench 82.7% (+13pts vs Opus 4.7)，自主 Agent 维度建立领先。Token 效率显著提升（第三方估算约 +40%），部分抵消 output 翻倍涨价。

82.7
Terminal-Bench

58.6
SWE-Pro

\$5/\$30
API \$/M

48天
距GPT-5.4

OpenAI 博客 / Artificial Analysis

长上下文跃升：1M Graphwalks BFS 9.4% → **45.4%**，Parents 44.4% → **58.5%**。数学发现：发现 Ramsey 数新渐近性质证明，经 Lean 验证。**安全**：CyberGym 81.8%，首次获"High"网络安全评级。

内部验证：85%+ 员工周活使用 Codex；财务团队审核 **24,771 份 K-1 税表**，加速 2 周。外部评价：NVIDIA 工程师称"失去 GPT-5.5 像肢体被截肢"，Cursor CEO 称复杂长时间任务显著更强。

高 高置信度 MIT 开源

DeepSeek V4
DeepSeek | 4/24 | V4-Pro + V4-Flash, MIT 开源

1.6T/49B MoE, 1M context, MIT 开源——开源阵营最大赌注。性价比革命：\$3.48/M output（标价，当前 75% 限时折扣至 5/5，实际 \$0.87/M）= Opus 4.7 的 1/7。竞赛编程之王：LiveCodeBench 93.5 超越所有模型。

55.4
SWE-Pro

93.5
LiveCodeBench

\$1.74/\$3.48
V4-Pro

\$0.14/\$0.28
V4-Flash

DeepSeek / Simon Willison

架构创新：Token-wise Compression + DSA（DeepSeek Sparse Attention），1M context 下 FLOPs 仅 V3.2 的 **27%**，KV cache 仅 **10%**。V4-Flash 更降至 10%/7%。HuggingFace 865GB——当前最大开源模型。**自评（公开模型卡）**："显著缩小了与领先闭源模型的差距"。

退役公告：V3.2 的 deepseek-chat 和 deepseek-reasoner **7/24** 后不可访问，迁移窗口 3 个月。

高 高置信度 MODIFIED MIT

Kimi K2.6
Moonshot AI | 4/20 | Agent Swarm, Modified MIT

1.1T/32B MoE（HuggingFace 确认 1.1T 总参），Agent Swarm 扩展到 300 并行子代理 + 4,000 协调步骤。SWE-Pro 58.6% = GPT-5.5——开源首次追平闭源。HLE (w/ tools) 54.0% 领先所有模型。

58.6
SWE-Pro

300
子代理

~\$0.60/\$2.50
API \$/M

54.0
HLE (w/ tools)

MarkTechPost

实验验证：Zig 推理优化（12h）：吞吐量 15→**193 tok/s**。金融引擎重构（13h）：**185%** 性能提升。5 天自主运维。

架构亮点：动态任务分解（搜索+研究+文档分析+写作+生成并行）/ Skills 能力 / Claw Groups（多代理+人类混合）/ 双模式（Thinking + Instant）。**商业限制**：Modified MIT 限制月活 >1 亿或月收入 >\$2,000 万的实体。

编程 Agent 不再是一维竞争。三条截然不同的架构路线已经确立——各自优化不同场景。多模型路由策略成为最优解。



编程修复

Claude Opus 4.7

64.3%

SWE-bench Pro

自验证，可靠修复。适用于真实 GitHub issue 修复。



自主 Agent

GPT-5.5

82.7%

Terminal-Bench 2.0

CLI 工作流，自主规划。领先第二名 13pts。



并行代理

Kimi K2.6

300

Sub-Agent Swarm

多代理协调，性价比高。已有生产验证。

3. Additionally, Google and XAI have started to urgently improve the coding capabilities of their models

- **Google DeepMind Forms "Strike Team" to Catch Up with Claude's Programming Capabilities (0421)**

According to The Information, Google DeepMind has assembled a "strike team" composed of researchers and engineers, focusing on enhancing Gemini's long-range programming capabilities, especially for complex multi-file code tasks. Sergey Brin's personal involvement directly reflects that Anthropic's Claude Code has already had a strong impact within Google—previously, Google engineers stated that Claude Code once completed the team's months-long workload in just one hour.

Insight by AI Claude Code is reshaping the "AI programming" standard, forcing Google to counterattack, and competition in the programming arena has entered a new stage.

[The Information](#) / [Heise](#)

- **SpaceX Collaborates with Cursor to Develop Programming AI, Secures \$60 Billion Acquisition Option (0422)**

SpaceX announced a partnership with Cursor to develop the next-generation programming AI, with the agreement including astonishing terms: SpaceX can acquire Cursor for \$60B later this year or pay \$10B as a cooperation fee. This combines SpaceX Colossus supercomputer (equivalent to a million H100 computing power) with Cursor's programmer distribution capabilities. Previously, two of Cursor's most senior engineers had jumped ship to xAI and reported directly to Musk, and xAI also began leasing tens of thousands of GPUs to Cursor. Cursor's valuation soared from \$2.5B in January 2025 to \$29.3B after the Series D round in November, and the acquisition price of \$60B was twice that of the Series D round.

Insight by AI Cursor still relies on Claude and GPT models to date, while Anthropic and OpenAI are launching their own programming tools to directly compete with Cursor.

TechCrunch

4. In the melee of world models, there is a special study to be updated (the same goes for AI4S; later, we will conduct a special study with Qianqian, one on AI4AI and the other on scientific base models)

In terms of embodiment, Sudu Pure Simulation is worth paying attention to, as there aren't many domestic players committed to this (if the problem of robots being resistant to falls cannot be resolved within two years, it still seems worthy of attention).

Synthetic data may dominate in two years. Real data acquisition may focus on types that are difficult to synthesize, such as tactile and force sensations (human fine manipulation), as well as whole-body coordinated data (human teleoperation) and interaction datasets (embodied agents truly enter the scene without fear of falling).

Signal 2: Strong demand for computing power + increasingly dispersed suppliers and geographical locations

First, the dispersion of computing power suppliers; second, the multi-location layout of data centers (on the one hand, the demand for computing power remains strong; on the other hand, such a layout is the key foundation for model companies to carry out global business)

|

分类	主体 / 主题	核心算力目标 / 核心结论	核心布局 / 核心特征	关键量化数据
头部厂商算力布局	OpenAI	2030 年总算力规模达 30GW，2028 年锁定 8GW 装机量	多元云厂商合作 + 自研芯片双轮驱动，分散供应链风险	1. 未来 5 年算力投入超 6000 单月算力运营成本超 10 亿美元 Stargate 项目 7 个数据中心在 2029 年预计达 9+GW 装机量
	Anthropic	2027 年前锁定至少 8.5GW 算力供给，支撑企业级业务扩张	AWS + 谷歌双云深度绑定，GPU+TPU 双路线产能锁定	1. 亚马逊追加 50 亿美元投资高 5GW 算力 2. 与谷歌合作 20 线 3.5GW TPU 集群 3. 每新增元年化营收，需配套 4GW 算力
	Google	2026 年新增超 5GW 算力，实现自研自用 + 对外商业化双轮驱动	全栈自研垂直整合，TPU 芯片 - 数据中心 - 云服务全链路自研	1. 发布第八代 TPU，史上首次训练专用 8t、推理专用 8i 芯片 2. 计划在印度建设美国外最大 AI 中心集群
行业供需	全球 DRAM 供需格局	供应缺口为 AI 驱动的结构性短缺，将长期持续	头部厂商产能向高端 AI 存储倾斜，通用产能收缩	1. 2027 年底全球 DRAM 产能 60% 市场需求，短缺将持续到 2030 年 2. 2026 年 AI 服务器已消耗全部的 DRAM 月产能
	AI 算力集群成本结构	AI 加速芯片是算力成本的核心刚性约束	行业融资核心流向算力硬件，中小厂商算力门槛持续抬高	GPU 集群全生命周期 TCO 中上融资额直接用于 AI 加速芯片
	中东算力需求	中东已成为全球算力市场核心增量变量	阿联酋通过资本 + 算力模式，快速构建本土 AI 产业能力	2025 年阿联酋 G42+MBZUA 了 Cerebras 86% 的营收
	Cerebras IPO 核心数据	晶圆级 AI 芯片已实现规模化商业化盈利	绑定头部大模型客户 + 中东市场，从硬件销售转向算力服务	1. 2025 财年营收 5.10 亿美元 (75.71%)，净利润 8790 万美 OpenAI 签订超 200 亿美元多协议
	CoreWeave 融资	第三方算力服务商成头部厂商核心算力补充，资本持续重仓	聚焦 AI 专用 GPU 集群与算力基础设施建设	完成近 160 亿美元融资，股价达 55%

Signal 3: Startups and Products Worth Attention

1. NeoCognition: Building a "Microcosm Model" to Achieve Agent Self-Learning (4/22&24)

NeoCognition, founded by Professor Su Yu from The Ohio State University (undergraduate from Tsinghua University, Ph.D. from UCSB, former Microsoft researcher), has secured \$40 million in seed-round financing, co-led by Cambium Capital and Walden Catalyst, with follow-on investments from Vista Equity Partners, Intel CEO Chen Liwu, and others. The team consists of approximately 15 members, most of whom are Ph.D.s. The core idea of the company is to enable AI agents to achieve continuous learning by constructing a "micro-world model" and transform

from generalists to domain experts. Currently, the success rate of AI agents is only about 50%, and NeoCognition aims to fundamentally address the unreliability issue of general agents.

2. **Flipbook: Say goodbye to traditional UI, and Agent interaction enters the native era of "real-time pixel rendering"**

Demo Video: <https://x.com/i/status/2046982383430496444>

Zain Shah (former OpenAI, YC S13) and two others released the prototype project Flipbook - imagine every pixel on the screen being rendered in real-time and streamed by a model, with no HTML, no layout engine, no code, and only presenting what the user wants to see. The tweet received 4.3M views and 22K retweets.

The transition from "model-generated code" to "model directly rendering pixels" represents a leap in interaction paradigms, similar to the shift from the command line to the GUI. Currently in the prototype stage, this concept has significant influence. If this path proves viable, it will fundamentally redefine the concept of "User Interface" - UI will no longer be designed but rather generated in real-time by the model.

- **Concept: Each pixel is real-time video**
 - No layout engine = automatic deformation + interactive anywhere
 - Technology Stack: LTX Studio Video Model + Modal GPU + WebSocket
 - Core Issue: Computing Power Cost

3. **[Generally Not Required to Read] Crunchbase's Top 10 Financing Deals This Week**

This week, only half of the top 10 financing rounds exceeded the \$100 million mark. In this era of booming venture capital, this situation is somewhat unusual. Nevertheless, some large-scale investments were approved, with Amazon's \$5 billion investment in Anthropic and the cooperation agreement leading the way. Other substantial financing rounds flowed to enterprises in fields such as aviation automation, vision therapy, and artificial intelligence analysis.

- Anthropic, \$5 billion, Foundation AI: AI giant Anthropic announced that Amazon will invest \$5 billion in the company, with the potential for an additional up to \$20 billion in funding in the future. Previously, Amazon had already invested \$8 billion in the San Francisco-based company. The latest financing also includes a partnership with Amazon for training and deploying Anthropic's AI assistant Claude.

- **Reliable Robotics, \$160 million, Autonomous Aircraft:** Reliable Robotics, a developer of autonomous aircraft systems, raised \$160 million in a new round of financing led by Nimble Partners. The 9-year-old company, with its headquarters in Mountain View, California, promotes its technology for commercial and defense aviation markets.
- **Ray Therapeutics, \$125 million, Vision Therapy:** Headquartered in San Diego, Ray Therapeutics is a biotech startup focused on vision rehabilitation therapy. The company secured \$125 million in Series B financing led by Janus Henderson Investors. Founded in 2021, Ray Therapeutics has raised \$247 million in venture capital and grants to date, according to Crunchbase.
- **Omni, \$120 million, AI Analytics:** Omni, a company developing an AI-powered analytics platform, has secured \$120 million in Series C financing, led by Iconiq Growth. This financing has set a valuation of \$1.5 billion for the four-year-old company with headquarters in San Francisco.
- **Tortugas Neuroscience, \$106 million, Biotechnology:** Tortugas Neurosciences, a biotechnology startup focused on neuroscience with headquarters in Framingham, Massachusetts, has successfully raised \$106 million in Series A financing. The founding investor, Cure Ventures, co-led this financing round with The Column Group and AN Venture Partners.
- **AcuityMD, \$80 million, Medical Technology:** AcuityMD, an AI-driven data and research platform serving clients in the medical technology industry, secured \$80 million in Series C financing. StepStone Group led the financing for this 7-year-old company headquartered in Boston.
- **OpenAI, \$75 million, Basic AI:** Robinhood Ventures announced that it has purchased \$75 million worth of Common Stock in OpenAI, a company headquartered in San Francisco. These shares are held by Robinhood Ventures Fund I, a publicly traded fund that provides investors with exposure to a curated portfolio of private companies.
- **Orkes, \$60 million, Workflow Orchestration:** Orkes, a vendor developing an AI-driven software workflow orchestration platform, secured \$60 million in financing during its Series B funding round. AVP led the financing for this 5-year-old Silicon Valley startup.
- **Courier Health, \$50 million, Health Tech:** Courier Health is a company dedicated to developing tools aimed at improving the healthcare experience for patients with chronic diseases or rare diseases. The company has successfully raised \$50 million in Series B financing. Oak HC/FT led the financing for this New York-headquartered company.
- **Serif Biomedicines, \$50 million, Biotechnology:** Headquartered in Cambridge, Massachusetts, Serif Biomedicines is a biotechnology startup focused on developing modified DNA as a new class of drugs. The company received \$50 million in initial financing from Flagship Pioneering at its inception.

4. [Generally skippable] Cloud Next 2026 Highlighted Startups

- **Star Company:**

- **Lovable** is expanding its use of Google Cloud by launching a new coding agent through Google's enterprise app marketplace. Lovable, a fast-growing ambient coding startup, had reached \$400 million in ARR as of February.
- **Notion** is the most popular document productivity application integrating artificial intelligence capabilities in Silicon Valley, with its latest valuation approximately \$11 billion. The application is leveraging the Gemini model to power its text and image generation capabilities.
- **Gamma** is an AI-driven PowerPoint killer app whose latest valuation has reached \$2.1 billion. Gamma is leveraging Google's cutting-edge image model Nano Banana 2 and other features of Google Cloud Computing Platform.
- **Inferact** is a **business reasoning startup** based on the open-source project vLLM. In addition to leveraging the tech giant's AI stack, the company also accesses Nvidia's GPUs through Google Cloud.
- **ComfyUI** is a popular open-source tool for creating AI-generated images and multimedia content. The tool also provides access to Nano Banana 2 and is leveraging other cloud capabilities.

- **Other startups:**



- Chorus Corporation has developed an intelligent label based on **artificial intelligence technology**, which can track the status and movement of goods in real time.
- Emergent AI, an atmosphere encoding platform.
- ExaCare AI, a company that develops artificial intelligence software for **post-acute medical care facilities**.
- Insilica is responsible for generating AI-generated regulatory-compliant **chemical safety reports**.
- Parallel develops technologies for building web search and research APIs for artificial intelligence agents.
- The artificial intelligence software developed by Proximal Health can **automate the insurance claim review process**.
- Reducto, a document parsing tool based on artificial intelligence technology.
- Stord, responsible for handling e-commerce fulfillment and parcel operations.

- Stylitics specializes in developing artificial intelligence image generation software for retailers, which is used to handle tasks such as clothing styling and product assortment.
- Temporal, a developer cloud environment built to prevent failures.
- Vapi, a development tool developed for building conversational voice agents.
- Vurvey Labs conducts **synthetic market research** through AI agents.
- Wand, an in-game assistant for single-player PC games.
- Watershed, a company that develops software to assist enterprises in reporting and managing sustainability programs.
- ZenBusiness is a comprehensive back-office work tool designed specifically for small businesses, which includes an AI chat assistant.

Signal 4: The blue ocean area that can be breached is precisely the area with high AI exposure and low perceived substitution risk

Anthropic Releases Survey on the Economic Impact of AI on 81,000 People

- Those engaged in **roles more exposed to AI** have more concerns about AI-driven job losses. These concerns are also higher among **early-career** respondents.
- **Workers in the highest and lowest income occupations reported the greatest productivity growth**, most commonly due to an expansion of the scope of work (taking on new tasks).




Speculative Sustainable Amendment:

- **High exposure x High perceived substitution risk** - The current hot area and also the area where the model is the product.
- **Typical:** Programmer, Web/Software Development


The earlier one is in their career, the more they worry about being replaced.

- For entrepreneurs, it is difficult to resist the direct encroachment of model bases unless they have extremely strong vertical domain data or end-to-end workflow optimization.
- Key Features:** The output is **highly digital and verifiable**.
- High exposure x Low perceived substitution risk** -- Opportunity area for developing Copilot (**indicating that user demand is high, but current model/product companies are not yet ready** -- if we believe that most tasks will ultimately move towards automation)
 - Typical:** Data Scientist, Financial Analyst, Legal Counsel, Market Research.
 - Why is the perceived risk low?** Because the core values of these occupations lie in “**Responsibility Taking**” and “**Complex Decision-Making**”. High AI exposure means it can handle 80% of basic tasks (data cleaning, case finding, draft writing), but the remaining 20% (causal inference, compliance risk, business judgment) is not yet well done by current models.
 - Opportunity: Deep tools embedded in professional workflows — harnesses and interaction layers that better understand business practices**. Users do not feel they will be replaced, so they **are willing to pay for increased efficiency** (willingness to pay is more rational and sustainable than that of the anxiety-driven group).



Workers in high-paying jobs, such as software developers, have gained the greatest productivity boost from artificial intelligence.

Managerial occupations have the highest productivity improvement, followed by development occupations, and finally research occupations and occupations that require some form of intervention or interaction.




At this stage, AI is still a capability amplifier for professionals: most people believe that the benefits of AI productivity belong to themselves (trying new tasks, improving efficiency), which also forms a verification relationship with the following figure.

10% of respondents indicated that employers or clients demanded and received more work. A smaller proportion mentioned the benefits to AI companies, and an even smaller proportion believed that AI would have a net negative impact. This depends on career stage: only 60% of early-career workers said

- **Low exposure, low perceived substitution risk** - either no opportunity or look at niche market
 - **典型:** Clergy、Chief Executives、Civil Engineers
 - **Opportunities are tricky:** They do not require a high level of intelligence, but AI cannot replace the role of humans because they involve a large amount of interpersonal communication, complex decision-making, or interaction with the real physical world, and the ROI does not add up. Unless there is a qualitative change in Embodied AI, there is indeed no entry point for pure AI products in the short term.
- **Low exposure x High perceived substitution risk** - Existing products may already be doing well; for example, chatbots can handle it, and parts involving judgment and taste may be considered more likely to be internalized by existing models/products
 - **Typical:** Photographer, Industrial Designer, Editor.
 - **Paradox:** According to Anthropic's algorithm, the logical exposure of these occupations may not be high, but the fear among practitioners is extremely high. This indicates that the impact of generative AI on aesthetics and creativity is subversive.

they personally benefited from AI, while the figure was 80% for senior professionals.



Currently, not many people feel that things are done better with AI; instead, more things can be done and the speed has increased

- **Opportunity:** Intent Engineering.


Enabling professionals to better express their ideas in a way that AI can understand and generate is the most advanced thing in this field. When AI can mass-produce works that score 80 points, the anxiety point for human practitioners lies in how to prove the remaining 20 points.

Epoch AI Survey: 80% of Claude Users Earn Over \$100,000 Annually

A joint survey by Epoch AI and Ipsos (three rounds with approximately 5,000 respondents) shows that 80% of Claude users come from households with an annual income of over \$100,000, while only 37% of Meta AI users reach this level. Approximately 50% of adults in the United States fall into this income bracket. The concentration of high-income users of Claude far exceeds that of other AI products, with the income distribution of users of other major manufacturers (Google, OpenAI, etc.) being relatively even, with 56-64% having an income of over \$100,000.

The user profiles of AI products are diverging. Claude targets professional/high-income groups (as a replacement for advanced productivity), while Meta AI covers the mass market (general entertainment). Different user bases are matched with their respective product evolution directions and monetization strategies.

Claude users skew towards higher-income households; Meta AI towards lower-income



Source: Epoch AI/Ipsos survey. Three waves: March 3–5; March 13–15; April 3–5, 2026. Waves are pooled to increase sample size across income groups. Estimates are weighted to be representative of the adult US population.

Overview of Key Information This Week

	Capacity Leap	Efficiency Revolution	Market Dynamics
Layer 0: Energy Layer	/	/	[Computing Power Tracking] 0420 Epoch AI is investigating Stargate, with all 7 Data center sites under construction, expected to reach 9+ GW by 2029 0421 Amazon to Invest Additional \$5 Billion in Anthropic (Total Could Reach \$20 Billion), Anthropic to

			Receive Up to 5 Gigawatts of Computing Power 0424 AI Giants Accelerate Computing Power Layout in Europe: Anthropic's Self-Build + Google's First Entry into Austria 0426 Google plans to invest up to \$40 billion in Anthropic and sign a 5GW computing power agreement at the same time
Layer 1: Computational Paradigm	0420 Google and Marvell are in talks to jointly develop AI inference chips 0420 Samsung halts production of LPDDR4/LPDDR4X 0422 AMD and GlobalFoundries collaborate to develop MI500 CPO optical interconnection solutions 0423 Google released the 8th generation TPU: training chip 8t and inference chip 8i 0421 Positron shipped its first AI chip in 18 months and won an Oracle customer 0425 Meta signed a multi-million-dollar large order for Graviton cores for Agent inference workloads	0421 SemiAnalysis Conducts In-Depth Analysis of GPU Cluster TCO, with Over 80% of Financing Used for GPUs 0425 Google DeepMind Releases 0421 Decoupled DiLoCo: 20x Speedup in Cross-Data Center Training	
Layer 2: Model and Cognition	【Base Model】 0421 Alibaba released Qwen3.6-Max-Preview to preview the next-generation flagship (closed-source) 0421 Kimi K2.6 went live on OpenRouter, supporting	【Edge Side, Small Model】 0423 Shangtang Jueying released the Sage Edge MultiModal Machine Learning model, with 3B parameters outperforming GPT-5.4	【Evaluation】 0420 Berkeley RDI Audit of 8 Major AI Agent Benchmarks, All Can Be Scored Full Marks by Zero-Capability Agents 0423 NVIDIA and Google Cloud have reached a decade-


long-range Agent programming 0423 Xiaomi released MiMo-V2.5: Native MultiModal Machine Learning + Agentic capabilities, with a million-token context 0424 OpenAI released GPT-5.5, achieving comprehensive breakthroughs in programming and knowledge work 0424 DeepSeek V4 was released [World Model & Embodiment]0421 Sudo Technology released the embodied model Sudo R1, achieving a 98% grasping success rate in zero-shot scenarios (one of the few domestic embodied companies focusing on the simulation route) 0424 ByteDance released Seed3D 2.0, a Coarse-to-Fine two-stage 3D generation 0421 Independent Variable Robotics Releases WALL-B: The World's First Embodied Foundation Model with a Globally Unified Model Architecture 0421 QuatRoPE: Quaternion Rotation Position Encoding Breaks Through 3D Spatial Reasoning [AI4S] 0422 Google 发布 Gemini Deep Research Max	0425 xAI Releases Grok Voice Think Fast 1.0: Tops the Voice Agent Leaderboard 0422 Ant Group 发布 Ling-2.6-flash 0423 Alibaba Releases Qwen3.6-27B Open-Source Model, Positioned as Flagship-Level Programming Capability • Research Focus [Inference Optimization] 0424 FASTER: Modeling the denoising process as an MDP significantly reduces the computational cost of diffusion strategies 0424 Poly-EPO: Exploring Reasoning Model through Set RL Training 0425 Fudan University et al. Propose HERMES Framework: Accelerating Streaming Video Understanding by 10 Times 0420 Luma proposes two routes: the federated model and the mega model for AI systems	long partnership to advance agents and physical AI
[Agent infra]	/	

Layer 3: System and Platform	0422 OpenAI Codex launched Chronicle screen memory feature 0423 OpenAI released ChatGPT workspace agent feature 0422 X discussed that OpenAI Codex's users increased from 3 million to 4 million within two weeks, and AI programming tools entered a period of explosive growth 0424 Claude Multiple Product Updates: Agent Memory Function Enters Public Beta , External Application Connectivity Expanded to Life Scenarios		
Layer 4: Application and Agent Economy	【Agent】 0421 Claude supports the construction of real-time Data Dashboards and trackers in Cowork0422 SpaceX collaborates with Cursor to develop programming AI and secures a \$60 billion acquisition option 0425 NVIDIA Deploys Codex Across the Company 0426 AI cybersecurity company Rilian secures \$17.5 million in seed round financing 【Embodied】 0420 AutoNavi releases full stack embodied technology system ABot and intelligent guide dog Tutu	[PMF Verification] 0420 OpenAI acquires AI education technology company Chalkie, with over 500,000 teacher users globally 0420 Anthropic receives investment intention with an \$800 billion valuation 0422 NeoCognition Secures \$40 Million in Seed Round to Build Self-Learning Agents 0423 Cloud Next 2026: \$750 Million to Support AI Agents 0423 Epoch AI Survey: 80% of Claude Users Earn Over \$100,000 Annually 0423 Anthropic Releases AI Economic Impact Survey of 81,000 People; Higher AI Exposure Leads to	[Embodied Financing] 0421 Embodied Intelligence company Gemini Intelligence completed a Pre-A round financing of tens of millions of yuan. 0422 Independent Variable Robotics completed a nearly 2 billion yuan Series B financing, with four major internet companies gathering together.

Computing Power Tracking

This week, the total number of Openrouter token calls reached 22T, with a month-on-month increase of 5%, achieving the first month-on-month growth since half a month ago, but still lower than the peak before the consecutive negative growth. Kimi 2.6 1.58T, Claude sonnet 4.6 1.36T, DeepSeek V3.2 1.28T, Claude Opus 4.7 1.15T, Gemini 3 Flash Preview 1.04T

Top Models
Weekly usage of models across OpenRouter



LLM Leaderboard

This Week

Compare the most popular models on OpenRouter

1.	Kimi K2.6 by moonshotai	1.58T tokens new	6.	MiniMax M2.7 by minimax	790B tokens ↓18%
2.	Claude Sonnet 4.6 by anthropic	1.36T tokens ↓2%	7.	Grok 4.1 Fast by x-ai	747B tokens ↑30%
3.	DeepSeek V3.2 by deepseek	1.28T tokens ↓0%	8.	Step 3.5 Flash by stepfun	721B tokens ↑98%
4.	Claude Opus 4.7 by anthropic	1.15T tokens ↑279%	9.	Claude Opus 4.6 by anthropic	685B tokens ↓44%
5.	Gemini 3 Flash Preview by google	1.04T tokens ↓9%	10.	Gemini 2.5 Flash Lite by google	651B tokens ↑9%

Explore apps and agents [opting into usage tracking](#) by category

① Browse the new Apps & Agents directory to explore by category, discover trending tools, and see detailed usage stats.



1.		OpenClaw > OpenClaw is an open-source AI a...	2.24T tokens	6.		ISEKAI ZERO > AI adventures. Travel with your fa...	229B tokens
2.		Hermes Agent > Hermes Agent is an open-source,...	1.44T tokens	7.		Janitor AI > Janitor AI is a chatbot platform wh...	179B tokens
3.		Kilo Code > Kilo Code is an open-source AI c...	1.42T tokens	8.		Cline > Cline is an open-source AI coding...	146B tokens
4.		Claude Code > Claude Code is Anthropic's agenti...	817B tokens	9.		Lemonade > The AI tool for Roblox games.	122B tokens
5.		Descript > AI Video & Podcast Editor	297B tokens	10.		Roo Code > Roo Code is a VS Code extension ...	115B tokens

2026年4月13日

	Others	11.2T
	Claude Sonnet 4.6	1.38T
	DeepSeek V3.2	1.28T
	Claude Opus 4.6	1.22T
	MiMo-V2-Pro	1.15T
	Gemini 3 Flash Preview	1.14T
	MiniMax M2.5	1.05T
	MiniMax M2.7	961B
	Gemini 2.5 Flash Lite	595B
	Grok 4.1 Fast	573B
Total		20.6T

2026年4月20日

	Others	12.6T
	Kimi K2.6	1.58T
	Claude Sonnet 4.6	1.36T
	DeepSeek V3.2	1.28T
	Claude Opus 4.7	1.15T
	Gemini 3 Flash Preview	1.04T
	MiniMax M2.7	790B
	Grok 4.1 Fast	747B
	Step 3.5 Flash	721B
	Claude Opus 4.6	685B
Total		22T