[Weekly Briefing] Industry Narrative Before the Next Generation Breakthrough of the 0505-0517 Model (Linked with 0518-0524)

💥 - **Trend 1: ToB Deployment + Acceleration of Self-Evolution Exploration at L2-L3 Levels**
    - **One consensus:** The improvement in model capabilities brought about by Coding will approach its upper limit this year.
    - **Currently, the first key priority is to accelerate the deployment of ToB in real-world environments**, while gaining experience from large-scale real-world tasks, achieving agent self-evolution through experience scaling at both the application layer and agent layer, and extending the lifespan of the current paradigm.
    - **Meanwhile, the Frontier Research Institute will continue to seek breakthroughs in next-generation models.**
        - Self-evolution, short-term: systems engineering + post-training + feedback closed-loop; medium-term: new training methods are needed, especially experience-based learning, such as continual/online post-training; long-term: exploration of new algorithms + new architectures + new training paradigms, such as open-ended algorithms.
- **Trend 2: Interactive Experience Transformation**
    - Changes in the interaction layer are essentially paving the way for Agent execution: AI no longer merely answers questions but can participate in workflows with higher bandwidth, lower friction, and more continuous online presence.
    - A problem jointly addressed by PMF verification in the speech domain, Google AI Mouse, Coding Agent Console, etc.: how to enable AI to understand user intent, environmental context, and operation objects with lower friction.
        - **The real-time interaction model of Thinking Machines Lab** upgrades the interaction ability from cheating software harness to the model's native ability;

        - **Google AI Mouse** upgrades screen interaction from "click position" to "understanding objects".

- **Trend 3: Robots begin to verify reliable execution in convergent scenarios**
    - 2025 is more like a "performance year" for robots, with motion control issues basically resolved; 2026 may enter the "interaction year", with the focus shifting to human-machine interaction and autonomous decision-making in specific scenarios; 2027 is more likely to enter the "productivity tool year".
    - Cases such as Figure and Yushu demonstrate that robots are evolving from single demos to reliability verification after scenario convergence, for example, room tidying, meeting room cleaning, package sorting, etc.
    - In the longer term, robot competition is not just about the hardware of the robot itself, but rather a full stack closed loop of the robot itself + motion model + decision model + data infrastructure. Game data, simulation platforms, and real-world feedback will all become important assets for World Model Training.
- **Trend 4: Strengthening of the Model + Harness Narrative**
    - Developments such as Claude Code Agent View, OpenAI Codex Mobile, MiniMax Agent Teams, Pinecone Nexus, and AWS AgentCore Payments all point to the same trend: Agent competitiveness is increasingly directed towards sustainable, monitorable, authorized, and tradable task execution systems.
    - Artificial Analysis's Coding Agent Index also reinforces this point: the evaluation object is shifting from "model" to "agent + model + execution setting".**Harness has independent value, and Cursor remains competitive.**
- **Trend 5: Three Structural Signals of Monetization**
    - Ramp Data (50,000+ Enterprise Samples): Anthropic's B2B Share Surpasses OpenAI for the First Time: The "Safe + Reliable" Positioning is Eroding the First-Mover Advantage, and the B2B Market is Not a Winner-Takes-All
    - Keling AI Spin-off - Video Generation Establishes Independent Track
    - AI replacing white-collar workers is evolving from prediction to reality
- **Computing Power Tracking**
    - **Total volume maintains an upward trend:** Single-week growth is 1.2T, showing a slowdown compared to the previous two weeks
    - **Owl Alpha Surges Ahead**: Rushing into the Top 10 list this week as a free model, with a growth rate of 121%
    - **Hy3 Preview Free Model Stable Head**

- **Kimi K2.6 call volume dropped** by nearly 35%, **DeepSeek V4 Flash call volume increased** by nearly 1T
- **Key OP this week**
  - 【OP】Anthropic 与 OpenAI toB 战略解读
  - 【OP】AI 自进化方向与 Neo lab 梳理（持续更新）
  - 【OP】Thinking Machines Lab Interaction Models
  - 【OP】Google I/O 2026

Key Trend Insights

Trend 1: ToB Deployment + Acceleration of Self-Evolution Exploration at L2-L3 Levels

1. OpenAI and Anthropic's ToB deployment accelerates, with MiniMax following suit

1.1 Interpretation of OpenAI and Anthropic's ToB Strategies

【OP】Anthropic 与 OpenAI toB 战略解读

- **Continuing from the previous judgment**: OpenAI directly enters the productivity space (with deeper embedding) vs Anthropic focusing on refining end-user experience + opening distribution channels (developer mindset)
- **JV Differentiation Considerations**: OpenAI's DeployCo aims to "productize, internalize, and platformize" enterprise AI deployment capabilities; Anthropic's service JV aims to "capitalize, network, and service-provide" Claude's enterprise deployment capabilities.
- **Future Trends**: When the technologies supporting Agentic capabilities spread to the general public, OpenAI has a greater advantage; the key to OpenAI's long-term success is**cost efficiency** (such as self-built Stargate and in-depth cooperation with Cerebras, etc.).
  - Sam Altman admits that he would prefer models to be**cheaper and faster**rather than smarter, but**users still value capabilities most.**The reality of the current AI market - although inference cost and speed are pain points, the capability gap of models is the core decision-making factor for users' choices.
- **MiniMax seems to have slowed down the phased training of its base model**

💥 Extended Thinking

- From the Internet to the mobile Internet, in the early stages of industrial development, large-scale paid services have always been led by B-side enterprises first. This is a period of infrastructure construction --> commercialization and monetization. **At the current stage, the key to competition on the B-side is model capabilities, while on the C-side, it is still more about cost efficiency.**
- High-value combinations in the Internet era: "efficiency + innovation", "innovation + lock-in", "efficiency + complementarity + lock-in", and "efficiency + innovation + complementarity + lock-in". In the AI era, innovation is no longer a bonus but a fundamental requirement; the effectiveness of lock-in also remains to be verified (refer to the one-click memory migration function and the development of harness). The value that "innovation + lock-in" can create is questionable.
- **The opportunity in the Harness domain may lie in how to provide a model-agnostic framework to help users painlessly migrate the underlying call model** . From this perspective, the debate over whether the improvement of model capabilities will overshadow the Harness layer may not exist. Assuming that model capabilities still need to be continuously improved within the visible range, unless everyone has a model and the model can self-evolve (both in-context & weight updates)
- **A good ToC product should not follow the same logic as that of the mobile internet era. It should not attempt to profit by occupying users' attention, but rather help users reduce entropy.** Humanity needs to return to simplicity and seek the true self. Products that increase attention should instead train people or help them explore their self-interest.
    - A vital product must be one that can help human evolution. It is not about making people more addicted, but about making them more free.

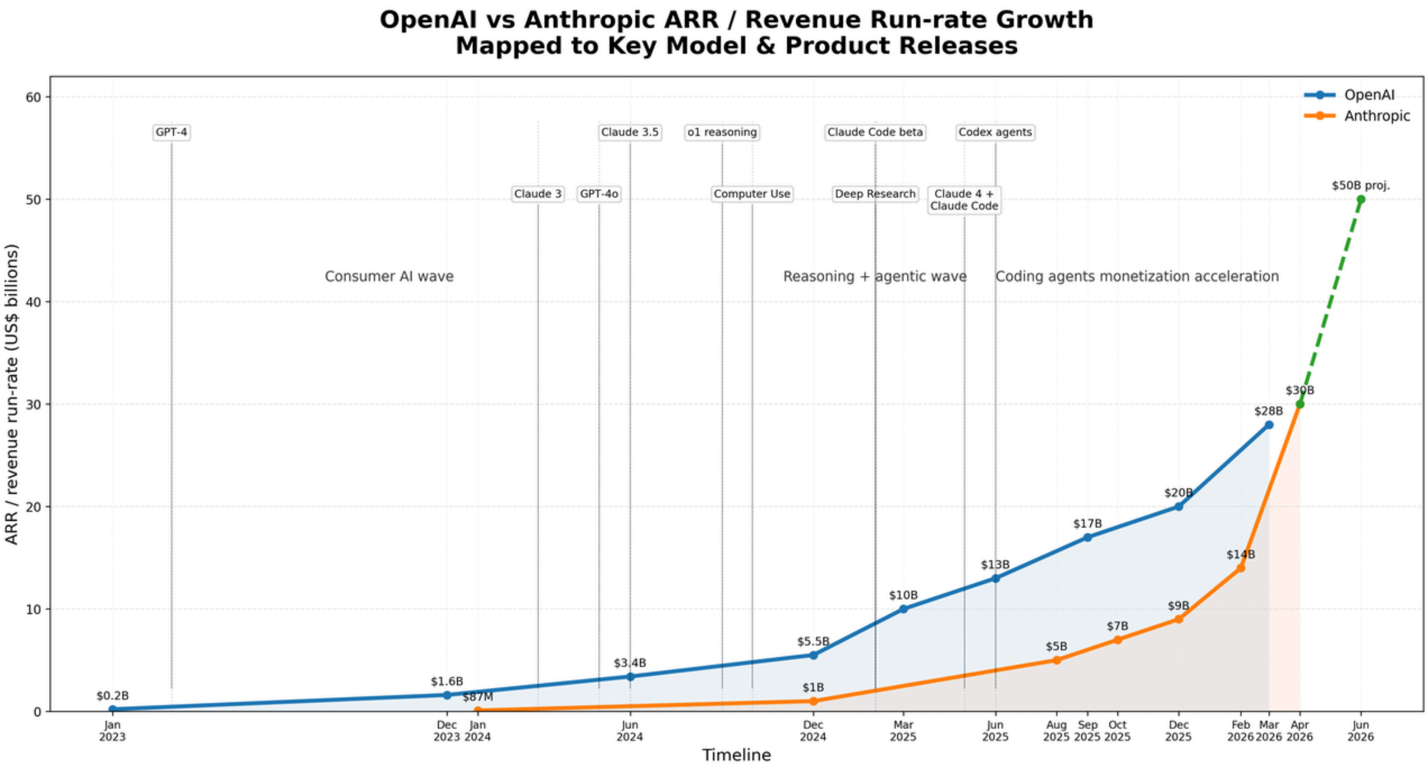

Historical ARR / revenue run-rate points are based on company disclosures and major media reporting (Reuters, WSJ, The Information). Dashed segments indicate forward projections / market expectations rather than audited GAAP revenue.

Differences between the two parties in the JV

Dimension	OpenAI Deployment Company	Anthropic Enterprise AI Services Firm	Core Differences
Organizational Nature	Majority-owned and controlled by OpenAI, similar to an extended deployment of OpenAI	Standalone entity, where Anthropic engineering and partner resources are embedded	OpenAI has stronger control; Anthropic has stronger independence
Strategic Intent	Integrate enterprise AI deployment capabilities into the closed loop of its own platform	Expand the enterprise service delivery radius of Claude	OpenAI pursues "platform closed-loop"; Anthropic pursues "ecosystem expansion"
Delivery Model	FDE delves deep into client organizations, reconfiguring key processes, connection models, and client data/tools/control systems. Acquired Tomoro, directly obtaining approximately 150 FDEs and deployment experts	Engineers from the new company + Anthropic Applied AI Team jointly customize the Claude system for clients	OpenAI is more like Palantir-style frontline deployment; Anthropic is more like an AI-native SI
control rights	OpenAI holds a majority equity stake and controls the company	Independently operated and jointly initiated by Anthropic, Blackstone, H&F, Goldman, etc.	OpenAI takes delivery as its own capability; Anthropic takes delivery as an ecosystem node
Customer Entrance	Large enterprises, complex organizations, critical business processes, existing OpenAI customers, and partner networks	Medium-sized enterprises, PE portfolio companies, and enterprises lacking AI engineering capabilities	OpenAI is delving deeper into the enterprise sector; Anthropic is expanding more towards mid-sized enterprises
Capital Partner Role	Provides capital, customer network, and transformation experience, but OpenAI still controls the game	Provide funding, portfolio distribution, customer sourcing, and enterprise service platform capabilities	OpenAI uses capital to amplify self-controlled delivery; Anthropic uses capital to build external distribution channels
Relationship with Consulting /SI	Both cooperation and potential competition: OpenAI is directly involved in deployment itself	More emphasis on supplementing the existing SI/partner network	OpenAI may compress some of the space of traditional consulting firms; Anthropic is more like empowering the consulting ecosystem

Data and Feedback Closed Loop	It is easier to directly capture customer requirements, deployment models, product feedback, and industry playbooks	Customer relationship and implementation experience are partly embedded in independent service companies and partner networks	OpenAI has a stronger learning closed-loop; Anthropic has a lighter expansion leverage
Meaning of Business Model	Model/API/ChatGPT/Codex + Deployment Service + Long-term Enterprise Workflow Lock-in	Claude/API/Claude Code + partner solution + revenue from service company deployment	OpenAI is more full stack; Anthropic is more partner-driven

1.2 Minimax10xTeam: The Layout of Domestic Model Companies in the Vertical Domain Ecosystem

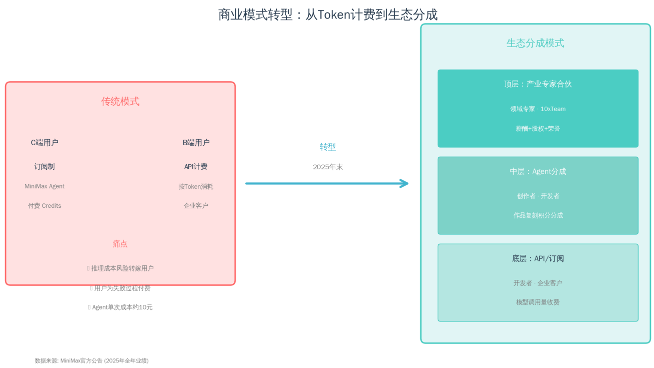

- Minimax 10xTeam Initiative: Recruitment is targeted at top experts in vertical fields such as industrial software, game engines, chip design, finance, and accounting, as well as fields globally that may be deeply integrated with large models, such as economics, life sciences, and materials chemistry. The initiative aims to **jointly build workflows and evaluation systems through experts**, achieving efficiency leaps in vertical fields.
    - **Flexible Collaboration:** The first is for full-time researchers, who can participate in and lead the implementation of research in their respective fields over the long term; the second is the Research Fellowship Program designed for senior experts who do not wish to leave academia, with at least 4 months of on-site collaboration, covering work locations in Shanghai, Beijing, Hong Kong, San Francisco, and London.

- **Exclusive Incentives:** MultiModal Machine Learning model capabilities, R&D environment, unlimited tokens, and collaborators can receive support such as stock incentives, joint academic co-authorship, and sharing of academic achievements.

- **Brainstorm: Does Minimax seem to have slowed down the competition in basic model iteration in stages (after all, the capabilities of domestic models still have some ground to catch up with compared to overseas top ones + IO once said that user data would not have a significant impact on model capabilities)?**
  - **Differences from Anthropic's STEM Fellow Program: Anthropic's program is a short-term research initiative involving on-site collaboration with domain experts.** It aims to leverage the professional judgment of top experts from various disciplines to **calibrate the Claude model and assist Anthropic in developing an "AI scientist" with autonomous long-range reasoning capabilities**. The program requires candidates to work full-time in its office for three months, with a weekly stipend of $3,800, approximately equivalent to 100,000 RMB ($14,720) per month. Selected candidates will gain access to the cutting-edge Claude model and internal evaluation tools, and be assigned an internal Anthropic researcher as a one-on-one mentor to collaborate on well-defined research projects.
  - **OpenAI Residency Talent Residency Program** targets professionals specializing in interdisciplinary fields such as physics, neuroscience, and mathematics, but who are passionate about artificial intelligence research. It aims to provide a pathway for outstanding talents with strong mathematical, scientific, or engineering backgrounds but who have not previously delved deeply into the field of artificial intelligence/Machine Learning to transition into top-notch **AI researchers or engineers**. There is an opportunity to be directly converted to a full-time position, becoming a full-time researcher or engineer at OpenAI. The **6-month** full-time program is based at the headquarters in San Francisco, US.Monthly salary is approximately **$18,300** (equivalent to an annual salary of approximately $210,000), and includes relocation support and full-time employee benefits.

1.3 Other relevant market analysis

Stripe Data Analysis: Revenue Growth Rate of AI Startups Significantly Faster Than SaaS

【OP】AI 公司收入数据分析|Stripe

- **The revenue growth rate of AI startups is significantly faster than that of SaaS** : The median time for the top 100 AI companies on the Stripe platform to reach $1 million in annualized revenue is only 11.5 months, about 4 months faster than the fastest group of SaaS companies.
- **New-generation AI companies monetize faster than early AI companies**: Among the top 100 AI companies, 55 were founded before 2020, and 45 were founded between 2020 and 2023. The median revenue generated by newer AI companies in their first year of operation is more than 7 times that of older AI companies in their first year.
- **AI companies have been global enterprises from day one:** The number of countries with sales for the top 100 AI companies in their first year on the Stripe platform is approximately twice that of SaaS companies, and this advantage persists until the fourth year.
- **Content generation AI products have the highest degree of internationalization**: In different AI usage scenarios, the median proportion of international revenue for content generation products reaches 75%, significantly higher than that of other usage scenarios, which generally fall below 52%.
- **Currently, AI companies still mainly focus on horizontal tools, but the trend towards verticalization is strengthening.** Among the top 100 AI companies in Stripe, 80 offer horizontal applications, such as content generation, chatbots, AI infrastructure, data synthesis, and process automation. However, the AI industry is following the path that SaaS once took, gradually moving from general-purpose tools to industry-specific solutions.

OpenAI B2B Signals Research: Leading Enterprises' AI Usage Intensity Reaches 3.5 Times That of Ordinary Enterprises

【OP】toB agent 渗透｜OpenAI研究0506

- OpenAI B2B Signals is a privacy-protected aggregated analysis platform based on enterprise usage data.
- **Key Findings**
  - **The AI competition gap among enterprises is shifting from whether they use AI to how deeply they use it.**
    The AI intelligence usage per employee in leading enterprises has reached **3.5 times that of ordinary enterprises (2 times a year ago)** . The volume of messages sent to AI can only explain **36%** of the leading edge, with most of the gap coming from richer and more complex AI usage.
  - **Agent workflow is becoming a cutting-edge symbol,**Agentic tools have the largest gap: the usage of Codex by leading enterprises is 16 times that of ordinary enterprises.

The greatest advantage lies in advanced tools, where the number of Codex messages sent by each employee at leading-edge companies is 16 times that of ordinary companies. Ordinary enterprises mostly use AI as a tool for Q&A, writing, and searching; leading-edge enterprises, on the other hand, have begun to use AI as a work agent capable of handling multi-step tasks. Cisco Case: Codex helps reduce build time by approximately 20%, saves over 1,500 engineering hours per month, and increases defect resolution throughput by 10-15 times. (0507)

- **Horizontal adoption of AI has already taken place, but the real value comes from vertical specialization, as different industries use AI in different ways.** Writing and communication remain the most common scenarios, but different functional departments are developing their own AI usage patterns: IT and security tend towards operating guidelines and process instructions, software development and data science lean towards coding, and finance leans towards analysis and calculation.

  The professional, scientific, and technical services industry leads in Codex and API usage; the finance and insurance industry leads in the widespread deployment of ChatGPT; the education services industry leads in per capita message intensity; and retail and healthcare also stand out in API usage intensity.

2. Accelerated self-evolution exploration at L2-L3 levels

AI Self-Evolution Direction and Neo Lab Review

【OP】AI 自进化方向与 Neo lab 梳理（持续更新）

1. **AI self-evolution is a new scaling direction from "scale model" to "scale experience closed-loop"** — a closed-loop where the model, agent, and environment co-evolve, one of the core paths leading to superintelligence.
2. **L3 (Automated Research Closed Loop) is the current verifiable inflection point, while L4 (Capability Feedback Flywheel) is the real watershed** , i.e., whether experimental trajectories, failure samples, environmental feedback, and research results can stably flow back to the next generation of models or Agents, forming a data/model/workflow flywheel
3. **Frontier model labs still has resource advantages to advance L4–L5**, considerations for startup opportunities:
   a. **Reasoning Self-Improvement (L1–L2):** Build reasoning harness, test-time optimization, and self-optimizing prompt/workflow/tool-use systems on top of existing models, similar to Poetiq. It can be productized quickly, but long-term barriers depend on whether cross-task reasoning experience can be accumulated and become the default control layer for model invocation or agent execution.

b. **New algorithms / infra / training paradigms (L4–L5):**Explore experience-based learning, continual / online post-training, and open-ended algorithms. The radical version, similar to Recursive Superintelligence, directly bets on recursive self-improvement;A more conservative version, similar to Sakana, starts from relatively verifiable directions such as evolutionary search, model merging, and automatic research closed-loop, gradually approaching the return of capabilities.

c. **Vertical Domain AI Scientist (L3–L4):** In high-feedback or high-value R&D scenarios such as chips, materials, drugs, robotics, and energy, construct a closed loop of hypothesis → simulation / experiment → feedback → next design. Similar to Recursive Intelligence in AI chip design or vertical AI Scientists in the materials / drug fields.The core barriers come from proprietary environments, simulation/experimental feedback, industry data, and design asset flywheels.

4. **The current main battlefield is to advance L2–L3 self-evolution in strong feedback scenarios.** Short-term industrialization opportunities mainly lie not in open recursion self-improvement, but in enabling Agents to continuously improve from cross-task experience, failure trajectories, environmental feedback, and experimental results. The real experience closed-loop of L2–L3 can accumulate infrastructure and Data assets for the future closed-loop of L4 capability backflow.

> **L2: Cross-task experience accumulation.**
> The system can reuse memory, tool usage patterns, workflow patterns, and failure experiences from past tasks in future tasks. Typical forms include enterprise agent memory, workflow optimizer, tool-use learning, and codebase/business system context accumulation.
>
> **L3: Automatic research/experiment closed-loop.**
> The system can autonomously complete the cycle of hypothesis → experiment → feedback → next hypothesis around a certain type of task. Typical forms include AI research automation, automatic experiments of coding agents, automatic parameter tuning of ML pipelines, AI Scientists, and autonomous experiments in simulation environments.

Review of Related Dynamics

- **Self-evolution Exploration of Frontier Model Companies (Examples from Public Information)**
  - **Automated Research/Automated Engineering Closed Loop in Verifiable Environments**, such as AlphaEvolve (0508), Qwen3.7-Max (0521);

Qwen3.7-Max	Qwen officially disclosed that it completed 1,158 tool calls and 432 kernel evaluations during approximately 35 hours of continuous autonomous

	execution; third-party reports claim that it achieved approximately 10x geometric mean acceleration.(Qwen Studio)
Google DeepMind / AlphaEvolve	AlphaEvolve is a Gemini-powered evolutionary coding agent that generates candidate algorithms via LLM, then validates, filters, and iterates them through an automatic evaluator, used for mathematics, algorithm discovery, and computational system optimization. (Google DeepMind)

- **Closed-loop of multi-agent scientific discovery**, such as Gemini for Science / Co-Scientist (0521);

Gemini for Science Toolkit	Three experimental tools and the Science Skills Toolkit: Hypothesis Generation (based on Co-Scientist, generating and validating hypotheses through multi-agent "idea tournaments", with clickable references), Computational Discovery (based on AlphaEvolve and ERA, generating and evaluating thousands of code variants in parallel to accelerate computational experiments), Literature Insights (based on NotebookLM, structuring literature search results into tables for comparative analysis).Simultaneously released the Science Skills Toolkit, integrating over 30 life science databases (UniProt, AlphaFold, AlphaGenome, etc.), which can reduce structural bioinformatics and genomic analysis from hours to minutes on the Antigravity platform (Google)

- **Self-optimization of evaluators, harnesses, workflows, and inference architectures**, such as MiniMax M2.7 (0318), ByteDance Agent-World Self-Evolution Training Framework (0506).

MiniMax M2.7	MiniMax officially states that M2.7's internal harness can automatically collect feedback, build internal evaluation sets, and continuously iterate on architecture, skills/MCP implementation, and memory mechanisms. (MiniMax)
ByteDance / Agent-World	Proposed the self-evolving training arena: Automatically explore real-world environment topics, databases, and executable tool ecosystems through Agentic Environment-Task Discovery, synthesize verifiable and difficulty-controllable tasks; then, through Continuous Self-Evolving Agent Training, combine multi-environment reinforcement learning and dynamic task generation to identify capability gaps and conduct targeted training. (arXiv)

- **AI Self-Training**. Anthropic's next-phase pre-training is underway, with AI training AI. Karpathy has joined Anthropic's pre-training team led by Nick Joseph. Karpathy will form a new sub-team focused on leveraging Claude itself to accelerate pre-training research (0519)

- **Startup Dynamics**
  - **Core Automation** aims to raise $500 million to $1 billion in financing, founded by the former vice president of research at OpenAI, to establish a new paradigm of continual learning (0509)
  - **Recursive Superintelligence** Raised $650 million in financing with a valuation of $4.65 billion (0514)
  - NVIDIA Collaborates with **Ineffable** to Build Reinforcement Learning Infrastructure (0514)

- **Other Related**
  - Harness Self-Evolution: Fudan, Peking University, etc. Propose AHE: Enabling Automatic Evolution of Coding Agent's Harness, Terminal-Bench Surpasses Human Design for the First Time (0502)
  - The University of Manchester and others jointly proposed TACO, a Plug and Play, training-free self-evolving compression framework that enables Terminal Agent to autonomously compress useless context (0508)
  - Jack Clark, co-founder of Anthropic, said in Import AI that he now believes there is a 60% probability that recursive self-improvement will be achieved by the end of 2028, meaning that AI systems will be able to automatically improve their own capabilities (0506)

Mr. Tang's latest views on self-evolution

- **Summary**

The three pillars of technology - Memory, Continuous Learning, and Self-Judgment - are being realized through engineering "tricks": 1M+ context window and RAG address memory, monthly and weekly model updates achieve pseudo-continuous learning, and Opus 4.7 has demonstrated early self-correction. Most critically, the endgame of self-training: Anthropic models may already possess the base line capabilities to write code, clean data, generate synthetic data, and self-train.

- **Final Outcome of Self-Evolution (Original Excerpt)**

The most difficult yet most promising path is self-evolution. The current wave is extremely powerful. I suspect that models like Claude may have already reached the baseline for self-training: **they can write their own code, clean their own data, generate synthetic data, and then train with this synthetic data**. This may "waste" some computational resources, but it saves the most precious resources: human effort and time. In the era of Machine Learning, speed is crucial. Rapid iteration is precisely the key factor contributing to the cognitive gap between leaders and followers. **It is rumored that the cluster with 2 million chips that Claude will use next year is very likely to be built precisely for this purpose: self-training of autonomous models** .

- Technical Overview

1M Context: Necessary base line.

Memory and Continual Learning: Prerequisites, most likely to be addressed first through "clever" engineering means.

Utilization Environment: Breakthrough Point.

Self-judgment: Turning point.

Fully autonomous training: the ultimate goal.

Trend 2: Interactive Experience Transformation

Human-Computer Interaction

Thinking Machines Lab Real-time Interaction Model: 276B MoE, Transforming Real-time Interaction from Cheating Software Harness into Native Model Capability

【OP】Thinking Machines Lab Interaction Models

- **Core Innovation**: Elevating real-time MultiModal Machine Learning interaction from product harness to model-native capabilities, using 200ms micro-turns, multi-stream input/output, and native training objectives, **attempting to scale "interaction capabilities" alongside "intelligent capabilities"**.
- **Technical Key Points and Effects**:
  - 200ms micro-turn continuous processing of audio/video/text concurrency + encoder-free early fusion (similar to Meta Chameleon's early-fusion concept, but TML focuses more on**real-time interaction**: both audio and images need to be processed within a <200ms window to serve immediate decision-making)

技术点	含义	
Time-aligned microturns	将连续交互切成约 200ms 的微轮次	模型可以近
Encoder-free early fusion	图像、音频等输入更早进入统一模型流，而不是先经过独立 encoder 抽特征	减少模态间 模

  - Double-layer architecture: Interaction Model (real-time perception) + Backend Model (deep reasoning)

  - Existing commercial real-time models (GPT-Realtime, Gemini Flash Live) are unable to meet its new benchmark (the model's ability to actively or restraint participate in interactions at the right time and in the right way)
- **More like the interaction layer of the AI era:** Connects foundation models, devices, visual environments, voice conversations, and human workflows. **The goal is higher bandwidth, lower friction, and more continuous online human-machine collaboration**.

This week's voice interaction-related dynamics

- Muse Spark's cross-platform push notifications cover WhatsApp, Instagram, and AI glasses, with new voice conversation and shopping features added. Meta leverages its 2 billion user base to drive the scaling of AI assistants. (0514)
- Wispr Flow Launches Hinglish Voice AI in India, Accelerating User Growth, and Plans to Roll Out Broader Dialect Support, Local Recruitment, and Lower Pricing. (0511)
- AI Voice Agent Platform Vapi Completes $50 Million Series B, Valued at $500 Million, 100% Adopted by Amazon Ring (0513)
  - Amazon Ring selected Vapi for investment after evaluating over 40 AI voice providers during last year's holiday season (**the field is crowded**)
- Voice AI company ElevenLabs has revealed a new list of investors, including well-known figures such as asset management giant BlackRock, actors Jamie Foxx, and Eva Longoria. ARR $350 million → $500 million (+43%) (0506)
- OpenAI has launched three brand-new real-time speech models in its API: GPT-Realtime-2 (the first speech model with GPT-5 level inference capabilities, with its context window expanded from 32K to 128K and supporting parallel tool calls), GPT-Realtime-Translate (supporting real-time translation from 70+ input languages to 13 output languages, priced at $0.034/minute), and GPT-Realtime-Whisper (streaming transcription, priced at $0.017/minute). The speech model has entered from simple Q&A into **the stage of reasoning while conversing**, with 128K context + tool invocation enabling speech agents to be truly usable in production environments. (0508)
- xAI releases Grok Voice Think Fast 1.0, a voice agent supporting complex workflows. In the voice field, it directly competes with OpenAI, with the ability to handle complex workflows being the key differentiator. (0509)

New Interaction Experience between Humans and Screens: AI Mouse

- **Google Unveils AI Mouse Prototype**. Core Capabilities: The cursor not only knows "where it points" but also understands "what it points to" (pixels → structured entities), allowing users to directly execute AI tasks by pointing + voice (e.g., "Fix this", "Move that"). Google has already begun integrating this capability into Chrome (pointing at web content to directly ask questions) and will soon launch the Magic Pointer feature on Googlebook laptops. (0516)

Trend 3: Robots begin to demonstrate the reliability of task execution in convergent scenarios, with the key breakthrough in 2026 lying in human-machine interaction

💥

- **Implementation Progress Views (from the research of Ren Xiaoyu of Dongyi Technology)**: 2025 is the "performance year", with motion control issues basically resolved; 2**026 is the "interaction year"**, with the focus shifting to human-machine interaction and autonomous decision-making capabilities in specific scenarios; 2027 will be the "productivity tool year", where robots truly transition from "being able to interact" to "being able to work", starting to create real value.
- **Scenario Convergence + Reliability Verification**
  - Figure 03: 2 robots complete room tidying in 2 minutes → 200 hours of continuous operation + 250,000 package sorting
  - Unitree G1: Completes the meeting room tidying task under human interference
- **Full Stack Route: Both the ontology and the model need to be developed together**
  - Genesis: Choose full stack self-developed from model to hardware
  - Figure: Terminate cooperation with OpenAI in February 2025 and shift to self-developed VLA model Helix
  - 1X: Self-developed World Model
  - Industrial chain hierarchical view (from the research of Dongyi Technology's Ren Xiaoyu): The first layer consists of complete machine enterprises that integrate decision-making models, motion models, and robot bodies (the three form a commercial closed loop); the second layer comprises platform companies that provide simulators, simulation tools, and data generation tools; the third layer consists of computing power providers such as training cards and inference cards infra manufacturers.
- **Data Infrastructure: The Value of Game Data to the World Model**

  - Origin Lab Completes $8 million Seed Round Financing (0514): Building the World's First AI Data Circulation Platform Focused on Game Data Assetization, Transforming Physical Simulation Data in Game Engines into World Model Training Data
  - Representative cases of training world models with game data
    - Decart: Trains Oasis with millions of hours of Minecraft gameplay and corresponding user behavior records
    - General Intuition: Based on Medal's large-scale high-energy gameplay clip data (2 billion videos produced annually)
    - DeepMind's early Genie trained an action-controllable world model using approximately 30,000 hours of internet 2D platform game videos, and subsequent Genie 2/3 further advanced towards 3D, real-time, and interactive environment generation.

Convergent Scenario + Reliability Verification

- Case presentation and logistics sorting LIVE of Figure F.03: First, it completed the verification of matching human basic efficiency and fundamental effectiveness under specific scenarios; second, it completed the verification of system reliability, namely the ability to operate continuously and stably in real environments and autonomously handle abnormal situations (navigate to the maintenance area or charging pile without human intervention).

时间节点	能力里程碑
5月9日	两台F.03机器人在2分钟内完成房间清洁和床铺整理
5月14日	Figure F.03机器人开始7×24小时不间断自主包裹分拣直播
5月21日	进入直播第9天，累计运行167小时，分拣209,000个包裹
5月22日	直播结束，连续无故障运行200小时，最终累计处理包裹249,560个

  - Related
    - Unitree G1
      - Released**Real-time Voice Control**function, allowing users to directly control robots to perform diverse actions through natural language(0520);
      - Unitree Technology demonstrated the multi-task fully autonomous operation capabilities of the G1 robot equipped with the WVLA 2.0 model and onboard vision and tactile sensing systems for cleaning a meeting room, with human interference (0526).

        - (recall) Zhiyuan Robot LWD Framework takes embodied intelligence from "offline training + deployment" to "deployment is training" (0502)
            - In the latest demonstration, the robot can skillfully cut fruits (pears, cucumbers) and put them into a blender to make drinks. This framework compresses the data flywheel from the traditional offline cycle of acquisition → training → deployment into real-time online learning. The larger the fleet size, the faster the data flywheel operates, representing a key paradigm shift for large-scale implementation.

Full stack narrative of ontology + model

- Genesis AI Releases GENE-26.5 (0507):Genesis AI Releases GENE-26.5 Full Stack Robot System Achieving Near-Human-Level Operation, with All Components Self-Developed, from Dexterous Hand Hardware (20 Active Degrees of Freedom), 1:1 Mapped Non-Invasive Data Gloves, 3 Millisecond Low-Latency Control Middleware, to GENE-26.5 Foundation Model and Genesis World Simulation Platform.The Demo covers delicate operations such as egg beating, pipetting, Rubik's Cube solving, and piano playing, all executed in real time by a single model, with most skills requiring less than 1 hour of training data. The company secured $105 million in seed financing, setting a record for seed financing in Silicon Valley's embodied intelligence sector.
    - Related
        - LiberAI Completes Nearly 500 Million Yuan in Financing (0516): Embodied intelligence company LiberAI (Liu Songming) has completed cumulative financing of nearly 500 million yuan, jointly invested by Sequoia China, ZhenFund, etc. The company has successfully run through the entire process of "UMI Hardware - Data Acquisition - Large Model Training", focusing on the technical paradigm of integrating human UMI data with the world model.
        - Xiaoyu's views on the division of labor in the industrial chain: Decision-making model (WAM/Embodied Brain)+ Motion model + Ontology, simulation tools, and data generation (training infrastructure), Computing power (inference/training infra).

        - Figure publicly confirmed in early April that it had terminated its partnership with OpenAI. The trigger for the split was a call from OpenAI revealing that it was internally exploring humanoid robots, which made Figure realize that it "was cultivating a competitor." After a year of cooperation, Figure believes that models should be continuously tested on physical hardware rather than in simulated environments; additionally, its internal team has already begun end-to-end research. [https://www.humanoidsdaily.com/news/i-fired-them-brett-adcock-on-the-openai-split-robot-self-repair-and-the-kid-safety-test]. After terminating the partnership, Figure has shifted to self-developing the VLA model Helix, developed by over 50 engineers specifically for humanoid robots.CEO Brett Adcock also said in an interview in May that his team had gained an edge in the AI model dedicated to robots, thus ending the cooperation[https://www.youtube.com/watch?v=99pOdGEGu6s].
        - 1X: Self-developed World Model 1XWM, Self-developed VLA Model Redwood
            - Regarding the OpenAI collaboration - 1X publicly used GPT-4o voice integration in NEO Beta to achieve real-time conversation capabilities.
            - 1X integrates modules such as built-in LLM, Audio Intelligence, Visual Intelligence, and Redwood AI in NEO for speech/audio context understanding, conversation, and home task interaction; in addition,

Data Infrastructure

- Origin Lab completed an $8 million seed round financing (0514): Building the world's first AI data circulation platform dedicated to the assetization of game data, it transforms physical simulation data from game engines into world model training data. The physical rules (gravity, collision, fluid, etc.) in game engines are highly similar to those in the real world, and can generate labeled physical interaction data at low cost and on a large scale.
- Butianshi Technology (0505): Founded by Bai Yuli, the former head of NIO's AI platform, it has received its first round of financing led by Sequoia China. It focuses on the construction of Embodied Data Infra, providing a one-stop engineering system covering Data Acquisition, processing, annotation, and training scheduling for robot Model Training. Currently, it is recruiting positions such as Embodied Algorithm System R&D Engineer and Robot Bottom Software R&D Engineer, and the team is in the early stage of formation.
    - Related
        - Decart Oasis attracted one million users within three days of its launch,confirming the monetizable value of the real-time interactive video generation world model in the gaming industry[https://fourweekmba.com/decarts-3-1b-business-model-they-made-ai-minecraft-that-broke-the-internet-then-turned-it-into-a-profitable-business-in-11-months/#Investment_Thesis].

- World Modeling is the new generation ofpre-training paradigm: Jim Fanasserts that 2026 will be the "Year of Physical AI Explosion", and world model pre-training must cover 3D motion, proprioception, and haptics.
- Lightspeed recently in the field of AI data infrastructure other layout :
  - Lead Investment Luel's $31.2 million Seed Round (Luel mobilized a global contributor network of over 500,000 people across 96 countries to collect MultiModal Machine Learning data, put each submission through a proprietary QA process, and provide an auditable dataset. The project takes only a few weeks from launch to delivery.)

Trend 4: Strengthening of the Model + Harness Narrative

- Last Week's Updates: Persistent Task Queue + Multi-Agent Console + Permissions/Payment/Memory/Device Takeover

Dimension	Dynamic
Front-end form: Agent Console (Console)	• **Claude Code Agent View** has integrated multi-backend sessions, status monitoring, and human intervention into a unified view; • **OpenAI Codex** on mobile allows users **to monitor, approve, and continue tasks across devices** , and even securely use the app after the Mac is locked.
Backend Architecture:**Multi-Agent + Harness**	• Artificial Analysis explicitly shifts the evaluation object from the model to agent + model + execution setting (depth of thought), indicating that the harness has already become a variable of capability; • MiniMax Agent Teams (Driven by Leader-Worker-Verifier Finite-State Machine) • **The knowledge layer is evolving from passive RAG to agentic reasoning/memory infrastructure** : Pinecone Nexus (passive retrieval → active reasoning, 90% token ↓) • **Related**：Chroma Context-1（20B 参数 agentic search model）；Weaviate Engram(managed memory service for agents)
Business Closed Loop: Tradable Agent Economy	• Notion claims that the team has built over 1 million Custom Agents; • AWS AgentCore Payments enables agents to independently access and pay for APIs, MCP servers, content, and other agents, thus starting to complete the "execution - payment - governance - observability" business infrastructure.

AA Coding Agent Index: Harness has independent value, while Cursor remains competitive

Artificial Analysis' First Systematic Evaluation of "Model × Harness" Combinations:**Programming Agent/Harness + Model + Inference Gear**

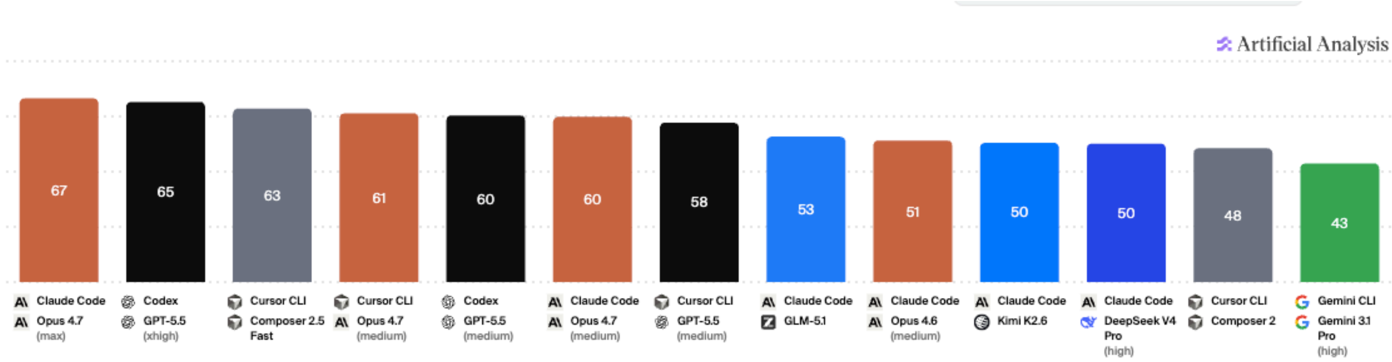

- **The value of Harness can exist, but the narrative of the integration of model + harness has not yet been broken**
  - Coding Agent Index：Opus 4.7（medium） + Cursor CLI = **61 分** vs Opus 4.7（medium） + Claude code = **60 分**
  - Opus 4.7 (medium) + Cursor CLI also performs better on SWE-Atlas-QnA
  - SWE-Bench-Pro-Hard-AA \ Terminal-Bench v2 shows little difference in performance

- **Cursor Remains Competitive: Cursor CLI + composer 2.5 Ranks Third**
  - **It is cheaper than all other agents in the index with scores exceeding 60 points:** Composer 2.5's single-task cost is only $0.07 (standard mode) and $0.44 (fast mode), placing it on the Pareto frontier of cost-effectiveness. The single-task cost of its medium-effort competing products at the same level ranges from $1.24 to $2.21; while the high-effort variants with scores 3 to 4 points higher cost as much as $4.10 to $4.82.
  - **Strategic Dynamics** : SpaceX collaborates with Cursor to develop programming AI (0422); Trains a larger model with 10 times the computational workload from scratch using million-level H100 equivalent computing power based on Colossus 2 (0520)
  - **Compared to Composer 2, Composer 2.5 represents a significant leap forward, primarily led by SWE-Bench-Pro-Hard-AA:** Composer 2.5's 47% score on SWE-Bench-Pro-Hard-AA is already comparable to Claude Opus 4.7 (max) in Claude Code.

- **Speed Leadership:** The average actual wall time for a single task in Composer 2.5's fast mode is 6.7 minutes, making it the third fastest agent in this index, trailing only Claude Opus 4.7 (medium) (5.8 minutes) in Claude Code and GPT-5.5 (medium) (6.2 minutes) in Cursor CLI.
- **Fast Mode offers better response speed at 6 times the price:** Fast Mode runs 30% faster than the standard version of Composer 2.5, but the cost per task is approximately 6 times that of the standard version ($0.44 vs $0.07). The Token pricing of Fast Mode is also 6 times higher: $3.00 / $15.00 per million input/output Tokens, compared to $0.50 / $2.50 for Standard Mode.

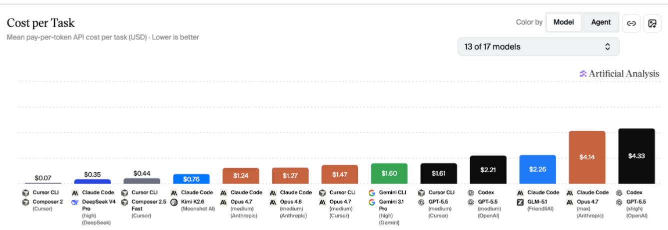

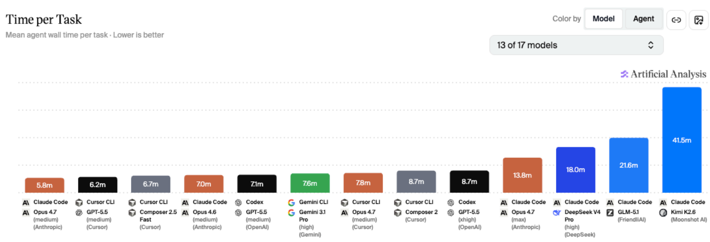

Other related benches: terminal-bench@2.0 Leaderboard

Rank	Agent	Model	Date	Agent Org	Model Org	Accuracy
1	vix	Claude Opus 4.7	2026-05-15	vix	Anthropic	**90.2%** ± 2.1
2	JJAgent	Multiple	2026-05-15	JJ	Multiple	**87.1%** ± 1.3
3	NexAU-AHE	GPT-5.5	2026-05-14	china-qijizhifeng	OpenAI	**84.7%** ± 2.1
4	LemonHarness	Multiple	2026-05-14	LR AILab of Lenovo CTO Org	Multiple	**84.5%** ± 2.6
5	Capy	GPT-5.5	2026-05-14	Capy	OpenAI	**83.1%** ± 2.1
6	Polaris	Multiple	2026-05-14	PolarisOps	Multiple	**82.2%** ± 2.8
7	Codex CLI	GPT-5.5	2026-04-23	OpenAI	OpenAI	**82.0%** ± 2.2
8	TongAgents	Gemini 3.1 Pro	2026-03-13	BIGAI	Google	**80.2%** ± 2.6
9	WOZCODE	Claude Opus 4.7	2026-05-14	WOZCODE	Anthropic	**80.2%** ± 2.1
10	LemonHarness	Multiple	2026-05-14	LR AILab of Lenovo CTO Org	Multiple	**79.9%** ± 3.0

Trend 5: **Three Structural Signals of Monetization**

Signal A: "Safe + Reliable" positioning is eroding the first-mover advantage. The B2B market is not a winner-takes-all market

- **Anthropic's B-side share surpasses OpenAI for the first time**
  - Ramp Data (50,000+ enterprise samples): **34.4%** of enterprises use Anthropic vs **32.3%** OpenAI

        - Within 12 months, Anthropic increased from 9% to 34.4%, while OpenAI decreased by 1% during the same period
        - Simultaneously launched Claude for Small Business (QuickBooks/Canva/DocuSign integration)
        - Valuation from $38 billion → negotiation **$90 billion** (nearly 3 times)

Signal B: Keling AI Spin-off - Video Generation Establishes Independent Track

- Kuaishou spins off Keling AI, valued at **$20 billion** (exceeding Kuaishou's total market capitalization of $29 billion)
- ARR**$500 million**(doubled compared to before the Spring Festival), 70% from professional user subscriptions
- Compared to Runway's valuation of $5.3 billion, Keling is nearly four times that amount

Judgment: Video generation AI has independently formed a capital-approved track. The $20 billion valuation indicates that the market believes this is not a subsidiary business of Kuaishou

Signal C: AI Replacing White-Collar Workers Transitions from Prediction to Reality

- **Cloudflare**: Lay off more than 1,100 employees (20%), CEO clearly states "AI is eliminating jobs", internal AI usage increased by more than 600% in March

Judgment: The evidence chain of AI replacing jobs has been upgraded from "companies announcing the use of AI" to "companies laying off employees due to AI + robots operating autonomously." This is happening simultaneously among white-collar workers (Cloudflare) and blue-collar workers (Figure)

Viewpoint Sharing & This Week's OP

【OP】Google I/O 2026

Key Points	Key Release	Core Meaning

The model enters the Agent Priority Era	Gemini 3.5 Flash released, Gemini 3.5 Pro previewed; Gemini Omni launched	The focus is not just on stronger models, but on a model stack that is cheaper, faster, and more suitable for performing actions and MultiModal Machine Learning generation. Gemini 3.5 Flash has been integrated into Antigravity, Gemini API, AI Studio, and Android Studio. (blog.google)
Antigravity 2.0 Becomes the Main Platform for Developer Agents	Antigravity 2.0 Desktop Application, CLI, SDK, IDE, Enterprise Edition	Google is transitioning from Gemini CLI to Antigravity, featuring multi-Agent orchestration, sub-Agents, background tasks, terminal sandboxing, credential masking, and Git security policies, directly competing with Claude Code / Codex / Cursor. (Google Developer Blog)
AI Studio has been upgraded from a prototype tool to a full stack App generation entry point	Vibe Coding、React/Angular/Next.js、Firestore、Firebase Auth、Secrets Manager、Android 构建与 Play Console 测试轨道	Google connects "prompt → prototype → backend → authentication → mobile testing and release → continued development of local Antigravity" into a closed loop, aiming to have AI Studio responsible for rapid exploration and Antigravity responsible for engineering implementation.Google I/O
Chrome is evolving into Agent Runtime	WebMCP、Chrome DevTools for Agents、Gemini Nano、Gemma 197M	WebMCP exposes structured tools to browser Agents, DevTools for Agents provides Agents with Console logs, network traffic, and accessibility trees, while the small models Gemini Nano / Gemma complement the edge-side execution capabilities. (Chrome for Developers)
Search / Shopping / Workspace 被 Agent 化	AI agents enter Search, Universal Cart, and Workspace API for native calls	Google's strength lies not in single-point agents, but in integrating agents into high-frequency entry points: scenarios such as search, shopping, work, and Gmail/YouTube, transforming tasks from "question-and-answer" to "background execution". (Reuters)
MultiModal Machine Learning creation continues to strengthen	Gemini Omni、Flow、YouTube Shorts Remix、SynthID	Gemini Omni is positioned as a model that generates multimodal outputs from arbitrary inputs, starting with video, and emphasizes physical understanding, editability, and SynthID watermarking. (blog.google)
Enterprise and monetization acceleration	AI Ultra price cut, new developer/enterprise pricing tiers, direct connection to	Google is pushing Agent into the enterprise production environment with lower-cost models, clearer pricing tiers, and a Cloud permissions

	Google Cloud, Antigravity Enterprise Edition	system. Reuters reported that Google is simultaneously launching cheaper enterprise models and adjusting subscription tiers. (Reuters)

Other

- 【OP】对推理SFT泛化性的重新思考
- 【研究底稿】从benchmark看模型前沿｜持续更新

Overview of Key Information This Week

	能力跃迁	效率革命
Layer 0：能源层	/	/
Layer 1：计算范式		0511 AMD ROCm性能14天内提升 DeepSeek V4发布助推 0510 AMD发布vLLM-ATOM插件 - GPU原生推理性能直通vLLM
Layer 2：模型与认知	**【基模】** 0506 OpenAI发布GPT-5.5 Instant - 替代GPT-5.3 Instant成为ChatGPT默认模型 0509 OpenAI发布GPT-5.5-Cyber限量预览版 - 面向经审核的网络安全团队开放 0515 AntLingAGI开源Ring-2.6-1T万亿参数模型 - Agent执行和推理双突破 0509 阿里发布CyberSecQwen-4B网络安全专用模型 0514 Claude Code上线Fast Mode - Opus 4.6速度提升2.5倍 **【世界模型 & 具身】** 0507 Genesis AI发布GENE-26.5全栈机器人系统 - 实现接近人类水平操作 **【AI4S】** 0510 谷歌DeepMind发布AI联合数学家 - FrontierMath Tier 4达48%创AI最高分 **【多模态】** 0508 OpenAI API上线新一代语音模型 - 支持推理、翻译和转录 0514 MiniCPM-V 4.6实现2.4倍推理吞吐提升 0508 xAI API推出图像生成质量模式 - 已在Grok生成3亿张图像 0511 Wispr Flow在印度推出Hinglish语音AI - 用户增长加速 0509 xAI发布Grok Voice Think Fast 1.0 - 支持复杂工作流的语音代理 0513 OpenRouter上线Perceptron-Mk1视频+具身推理VLM - 支持时间推理/密集计数/多摄像头联合推理	**【模型架构创新】** 0507 SubQ发布SSA架构 - 百万tok 理加速52倍 0512 Thinking Machines Lab提出 Model架构 - 276B MoE原生支持实 **【推理优化】** 0511 Databricks发布Genie数据Ag - 多LLM+并行推理，准确率32%→
Layer 3：系统与平台	**【Agent infra】** 0512 AWS Bedrock AgentCore支持AI代理自主支付 - 推Agent Toolkit简化开发	**【Agent infra】** 0517 Pinecone发布Nexus知识引擎 agent高达90%的token使用
	【Agent】 0515 xAI推出Grok Build早期测试版智能体CLI 0509 Anthropic正式发布Claude for Microsoft 365 - 覆盖Excel/PowerPoint/Word/Outlook 0512 Claude Code推出Agent view - 支持多会话并行管理与后台运行	

Layer 4：应用与智能体经济	0512 Claude Code推出Agent View - 支持多会话并行管理与后台运行 0507 Anthropic正开发Orbit - 为Claude Cowork提供主动式助理功能 0508 Google DeepMind发布AlphaEvolve进展 - 从研究走向实际问题解决 0513 Google发布Gemini Intelligence - Android从OS进化为智能系统，支持多步任务自动化和自然语言生成Widget 0506 Anthropic发布金融服务业Claude代理模板 【具身】 Figure 2台F.03机器人2分钟内完成房间清洁和床铺整理 【硬件】 0506 Apple计划在iOS 27中允许用户选择第三方AI模型驱动设备端功能 0517 开发者将M5Stack Cardputer改造为手持Claude设备 【安全】 0514 OpenAI为Windows版Codex构建安全沙箱	**【PMF 验证】** 0514 Anthropic推出面向小型企业 数据首次显示企业客户数超OpenA 0506 Microsoft Copilot Cowork扩 新增Skills和插件支持 0515 OpenAI在ChatGPT移动端推 预览版

Top 10 US Financing Amounts This Week

序号	公司	融资金额	领域	总部	轮次/估值	领投/投资方
1	Anduril Industries	50亿美元	国防科技	加州科斯塔梅萨	H轮；估值610亿美元	Andreessen Horowi Capital
2	VoltaGrid	7.75亿美元	能源	休斯顿	战略投资；另含2.25亿美元二级购买	Halliburton、Blacks
3	Mind Robotics	4亿美元	机器人	加州帕洛阿尔托	新融资；累计融资超10亿美元	Kleiner Perkins
4	Cowboy Space	2.75亿美元	太空技术	加州圣卡洛斯	B轮；估值20亿美元	Index Ventures
5	Oishii	1.5亿美元	室内农业	新泽西州泽西市	C轮；累计融资3.7亿美元	SPARX资产管理
6	Exa force	1.25亿美元	网络安全	圣何塞	B轮	HarbourVest、Peak Mayfield、Khosla Ve Seligman Ventures
7	Create Medicines	1.22亿美元	生物技术	未提及	B轮	Newpath Partners、Venture Partners、 Venture Partners
8	Havoc AI	1亿美元	自主系统	罗德岛州普罗维登斯	A轮；累计融资2亿美元	未提及
9	Star Catcher	6500万美元	太空技术	杰克逊维尔	A轮	B Capital、Shield C Cerberus Ventures
10	GridCare	6400万美元	数据中心电力	加州红木城	A轮	Sutter Hill Ventures

Computing Power Tracking

- **Total volume maintains an upward trend: growing by 1.2T per week, which has slowed down compared to the previous two weeks.**
- **Owl Alpha has emerged as a dark horse**: As a free model, it has entered this week's Top 10 list, with a growth rate of 121%.
- **Hy3 Preview Free Model Stable Head:** The rise and fall of the two-week call volume is extremely small, and the situation is stable.
- **Kimi K2.6 has declined, and DeepSeek V4 Flash has taken its place** : The call volume of Kimi K2.6 has dropped by 0.56T, a decline of nearly 35%; meanwhile, the call volume of DeepSeek V4 Flash has increased by nearly 1T.

2026年5月4日

Others	14T
Hy3 preview	2.68T
Kimi K2.6	1.61T
Claude Sonnet 4.6	1.45T
Claude Opus 4.7	1.24T
DeepSeek V4 Flash	1.11T
Gemini 3 Flash Preview	1.07T
DeepSeek V3.2	868B
Hy3 preview	857B
DeepSeek V4 Pro	816B
Total	25.7T

2026年5月11日

Others	14.1T
Hy3 preview	2.66T
DeepSeek V4 Flash	2.06T
Claude Sonnet 4.6	1.55T
Claude Opus 4.7	1.54T
Gemini 3 Flash Preview	1.15T
Kimi K2.6	1.05T
DeepSeek V3.2	1.03T
Owl Alpha	895B
DeepSeek V4 Pro	893B
Total	26.9T

Epoch AI Briefing: Servers Account for 60% of 1GW Data Center Costs, Claude Excels in SWE but Lags in Math

- Epoch AI released its May 15 briefing, with three key findings: (1) GovAI researchers modeled the total cost of a 1GW US AI Data center, where servers accounted for **60%** (**$5 billion** / $8.5 billion) of the annual total cost, far exceeding energy and other operating costs; (2) The Claude series of models, relative to its general ECI value, **exceeded the benchmark** on the software engineering benchmark and **fell below the benchmark** (based on the Domain-specific ECI Explorer); (3) The salary of AI super researchers reaches **10 - 100 times** that of their peers. Even a small difference in ability can lead to a significant salary gap. Additionally, Epoch conducted AI-assisted reviews of FrontierMath Tier 1 - 4 and found that approximately **1/3 of the questions had fatal errors** . Revised scores will be released after manual verification.

> 💡 A 60% share of server costs means that the leverage effect of improving chip efficiency on the economics of Data centers is far greater than that of energy optimization; Claude's bias towards SWE/mathematics provides quantitative basis for model selection

- Source: Epoch AI

Other interesting things

- Combination of software and hardware: The essence of C-end consumption - finding the**interaction form of current hardware**and**requirements**and value

Spillover of Coding Capabilities: Lovable Invests in AI-Generated Hardware Enterprises

Lovable Invests in Atech - Introducing Vibe Coding to Hardware Prototype Development (0515)

- **AI Application Building Platform Lovable led the investment in Danish hardware startup** Atech's $800,000 pre-seed round, followed by Sequoia Scout Fund, a16z Scout Fund, Nordic Makers, etc. Atech's product model: Users purchase a hardware starter kit, describe hardware concepts through AI conversations on the website, and AI automatically generates code to help build a working prototype. The user scope ranges from "4-year-old kids building cars to precise voltage sensing in hydrogen synthesis plants."
- **Business Model:**Atech has proposed the concept of "vibe-engineering" - users**describe hardware concepts through natural language, and the platform automatically completes underlying tasks such as circuit selection, firmware generation, and control logic writing. After purchasing the entry-level hardware kit, users can assemble it to obtain a working prototype**. Gustav Hugod, the head of customer experience, said that the user base ranges from "4-year-old children making toy cars" to hydrogen energy plants requiring precise voltage sensing.

- **Core Concept:**Atech CEO Tomas Harmer (Nordic Serial Entrepreneur in hardware) pointed out: "Software has a complete toolchain that allows a teenager to build an application in a weekend - hardware does not. Hardware currently remains at the first level of abstraction. Atech is building those missing abstraction layers." Lovable CEO Anton Osika said: "I see the pattern that Lovable once had, but in the hardware field. The team is very special."

About Lovable

> Users describe application requirements in natural language, and the platform automatically generates a runnable full stack web application, including front-end UI, back-end logic, database, authentication, API integration, and deployment infrastructure.

In July 2025, Lovable announced that its ARR had exceeded $100 million, stating that it took only 8 months to grow from its first $1 million ARR to $100 million ARR, and that more than 10 million projects had been created on the platform. In December 2025, Lovable announced the completion of a $330 million Series B financing, with a valuation of $6.6 billion. By March 2026, Business Insider reported that Lovable's ARR had grown from $300 million to $400 million, indicating that its growth continued. Lovable:

Other

In the field of AI-generated hardware, many 3D generation companies are involved, with more specialized tracks such as AI industrial design. Additionally, there are domestic companies similar to Atech, such as Wujie Ark ("Vibe Hardware" paradigm) and Tuya Cobuilder AI hardware zero-code generation.

创业方向	代表性公司/项目	核心切入点
AI+3D打印	Meshy（胡渊鸣创立，2025年ARR已年翻14倍，签约拓竹、xTool等头部3D打印厂商）	打通“AI创意→实体交付”全链条，字/图片即可生成3D模型并下单打印
AI+硬件原型生成	Schematik（被WIRED称为“硬件版Cursor”，已获Lightspeed 460万美元投资）	自然语言→ESP32/Arduino/Pico的（代码+接线图+组装指南+BOM），型验证
AI+工程软件革命	Orthogonal（前德国宇航中心团队创业，获36氪报道）	构建AI Native云端硬件开发平台，整计到制造的全流程，已有多物理场数解器

New Hardware (0506) at the Claude Developer Conference

Anthropic officially recommended its first official desktop companion intelligent hardware **Claude Desktop Buddy** at developer events such as Code with Claude. This hardware is essentially a physical AI terminal that collaborates with software, and its core hardware base uses modules manufactured in Shenzhen, China.

Hardware devices mainly have the following core highlights and functions:

- **Core Hardware Foundation**: The device is based on the development kit of Chinese hardware ecosystem manufacturer **M5Stack** (including `M5StickC Plus`, `StickS3`, and `Cardputer` and other hardware powered by the Espressif ESP32 chip).
- **Physical AI Interaction (Desktop Pet)**: Transforms the states in the Claude programming tool into physical feedback. The device can intuitively display Claude's current working status, turning the dull code review and approval process into a fun physical "desktop pet" interaction.
- **Physical Approval Control**: Through Bluetooth connection, when Claude Agent automatically generates code or executes actions, developers can directly approve or reject modification instructions proposed by AI with a single click via the screen and buttons on the physical hardware without switching desktops.

Developer Community Derivatives: After the release of Claude Desktop Buddy, a rich and diverse range of derivative works quickly emerged in the developer community, spanning productivity tools, parent-child education, gaming and entertainment, and creative peripherals—all supported by the same ESP32 hardware, which enables completely different interactive experiences in different scenarios. According to the official website of Espressif, developers have already created works such as the "Magic Wand" (shaking to switch between light and dark modes), a handheld marble maze motion-sensing game, a mini electronic keyboard, and the "Pocket Oregon" survival game.

> 💥 Extended thinking: Vibe coding has been well resolved on existing devices, so perhaps it is not wise to develop hardware around it? What exactly is a good **Agent Control Layer?**
>
> - Returning to the essence of the need, why do people need Vibe coding? What do people need when doing Vibe coding?
>   - **Entity AI Interaction (Desktop Pet) is emotional value**

Is space computing power just a buzz word?

> 💥 The Space Computing Power Industry Ecosystem Partner Program was established! The "Xingshu Plan" was launched simultaneously, with Shanghai seizing the new track
> https://mp.weixin.qq.com/s/C0RKF10mDxAb8Ku953I-Og?scene=1

Why go to space?

Energy, land constraints + support for distributed

Core Narrative:

At the beginning of 2026, the World Economic Forum in Davos became the turning point that triggered the topic of space computing power.

On January 22, 2026, Musk asserted at the Davos Forum that "space will be the location with the lowest cost for deploying AI, and this situation will become a reality within two years, perhaps three years, or at the latest within three years." He pointed out that the core of the logical viability of a space Data center lies in simultaneously having perfect conditions for both energy and heat dissipation - in space, there is not only an inexhaustible supply of solar energy, but also natural cooling can be achieved by simply placing the radiator on the backlit side, utilizing the extremely cold environment of deep space at approximately -270°C.

On January 31, 2026, SpaceX submitted an application for the "Orbital Data Center System" to the US Federal Communications Commission (FCC), proposing to deploy up to 1 million satellites in low Earth orbit to build a space orbital data center. The system will operate at altitudes ranging from 500 to 2000 kilometers, spanning multiple orbital layers to meet global demand.SpaceX pointed out in its application documents that the Orbital Data Center is the "most cost-effective and energy-efficient way" when AI computing demand grows rapidly.

In February 2026, SpaceX announced the acquisition of xAI, integrating rocket launches, satellite networks, and large model technologies, with plans to build a large-scale orbital Data center network. This move accelerated the industrial engineering process, forming an internal closed-loop ecosystem where SpaceX is responsible for launches and xAI provides requirements.

In March 2026, the US government introduced the <Protecting Taxpayer Commitment> plan, and companies such as OpenAI and xAI have signed commitments, clearly stating that they will independently build relevant infrastructure or cover all electricity costs, so that taxpayers will not have to foot the bill. This policy provides financial support for the construction of space-based Data centers.

By the end of March 2026, Musk first revealed the technical details of SpaceX's Orbital Data Center and made public the concept map of the "AI Sat Mini" satellite equipped with large radiators. The satellite will be equipped with heat sinks to dissipate heat in the space vacuum environment, with a designed operational lifespan of five years.

On May 6, 2026, Anthropic announced a partnership with SpaceX, securing computing power support from SpaceX's Colossus 1 Data Center located in Memphis, with a computing power scale exceeding 300MW (megawatts).More importantly, Anthropic has expressed interest in collaborating with SpaceX to develop "orbital AI computing power," with Anthropic "expressed interest" in partnering with SpaceX to develop multi-gigawatt orbital AI compute. This marks the transition of space computing power from technical verification to the stage of commercial cooperation, and the deep integration of AI model manufacturers and space infrastructure has become a reality

Recent Embodiment-related Financing

Company Name	Time	Financing Information	Core positioning/Technical path	Key data/Progress	Strategic Significance
Butianshi Technology	May 5th	In the first round of financing, Sequoia China led the investment	Embodied Data Infrastructure (Infra) provides a one-stop engineering system for collection, processing, annotation, and training scheduling	Founder Bai Yuli once led the NIO Edge-Cloud Computing Power System (378.1 EOPS, 15,000 concurrent nodes)	Transferring the "large-scale data closed-loop" methodology of intelligent driving to the embodied domain, with the goal of becoming the "Scale AI" of embodied intelligence. The key lies in whether we can accumulate general data processing capabilities across scenarios and tasks.
Guanglun Intelligence	May 25 (New Financing)	Valued at over $1 billion in March 2026 (the world's first unicorn in embodied	Simulated Synthetic Data - Batch Production of Robot Training Data Using High-Precision Simulation Technology	More than 80% of the simulation assets of major international embodied intelligence teams come from Guanglun; revenue will increase 10-fold in 2025, and new orders in Q1 2026 will reach 550 million yuan	Define the data and simulation infrastructure for the era of physical AI. However, it faces the Sim-to-Real gap: the success rate of manipulation in the simulation environment is 89.4%, but it plummets to 12% in real home scenarios. Simulation can scale up, but real closed-loop verification is irreplaceable.

		data); this round led by Ant Group			
Lingyu Intelligence	May 20th	Nearly 100 million yuan in Angel + Round	RaaS (Robotics as a Service) + MaaS (Model as a Service) dual-wheel drive, using high-performance ontology (TA robot, human arm isomorphic + three omni-directional wheel chassis) to collect real machine data	Intended orders are approximately 3 billion yuan, orders on hand are approximately 1 billion yuan, and 1,000 units are expected to be shipped this year	Generate revenue through the robot itself (leasing or selling), making Data Acquisition an endogenous part of the business closed-loop rather than relying on financing subsidies.
OriginFlow	May 21	Cumulative financing exceeds 500 million yuan	"NeuroScale" Paradigm: Non-invasive motor neural interfaces capture electromyographic signals, which are decoded and reconstructed by the base model PULSE to obtain data on hand posture, force exertion, tactile sensation, etc.	Founder Qin Shentao, a post-00s Tsinghua doctoral student	By bypassing the perspective occlusion and asynchronous errors of pure vision solutions, it is expected to become the "second-generation staple food" for embodied Data Acquisition. The ultimate vision: everyone contributes data unconsciously, achieving zero marginal cost supply.
Zhuma Innovation	May 25	Angel + Round	The world's first consumer-grade 3DGS (3D Gaussian	The first batch of consumer-grade devices is expected to be	Logic of "Dimensionality Reduction Acquisition": Make data acquisition devices consumer-grade products,

			Splatting) spatial camera, positioned as the entry point for real-world 3D data	shipped in the second half of 2026	collect naturally precipitated activity data (rather than laboratory-customized data), and form a large-scale real-world spatial data supply.
Origin Lab	May 14th	$8 million seed round	Game Data Assetization Platform: Transforming physical simulation data from game engines into world model training data	Focus on game data, providing low-cost, labeled, large-scale "synthetic physical world" data	Solve the "start-up problem" of data supply - before real data accumulates to a critical volume, use game-synthesized data to provide an immediately available supplementary solution.
Mifeng Technology	Zhiyuan Robotics Holdings has not disclosed the time of its independent financing	Zhiyuan Robot Holdings	The standardized and platformized output of Zhiyuan's self-built Data Acquisition Factory (over 4,000 square meters, with an average daily output of 30,000 - 50,000 data entries) provides general data services for the B-side.	Separation of Zhiyuan Ontology Business and Data Service Business	Upgrade embodied data from "self-used assets" to "marketable commodities". If a general supply capacity across ontologies and scenarios is established, its independent value may exceed that of the ontology business—transforming from a byproduct of gold mining into a water-selling business.